

# Subsetting Data in R

Data Wrangling in R

# Overview

We showed different ways to read data into R using:

```
readr::read_csv()  
readr::read_delim()  
readxl::read_excel()
```

In this module, we will show you how select rows and columns of datasets.

# Setup

We will be using the **dplyr** package in the tidyverse.

Here are several resources on how to use **dplyr**:

- <https://dplyr.tidyverse.org/>
- <https://r4ds.had.co.nz/>
- <https://cran.rstudio.com/web/packages/dplyr/vignettes/dplyr.html>
- <https://stat545.com/dplyr-intro.html>

The **dplyr** package also interfaces well with tibbles.

# Dataset

We will be using the `diamonds` dataset in the `ggplot2` package as an example (so make sure you initiate the `ggplot2` package if you are following along on your own).

```
head(diamonds)
```

```
# A tibble: 6 × 10
```

|   | carat | cut       | color | clarity | depth | table | price | x     | y     | z     |
|---|-------|-----------|-------|---------|-------|-------|-------|-------|-------|-------|
|   | <dbl> | <ord>     | <ord> | <ord>   | <dbl> | <dbl> | <int> | <dbl> | <dbl> | <dbl> |
| 1 | 0.23  | Ideal     | E     | SI2     | 61.5  | 55    | 326   | 3.95  | 3.98  | 2.43  |
| 2 | 0.21  | Premium   | E     | SI1     | 59.8  | 61    | 326   | 3.89  | 3.84  | 2.31  |
| 3 | 0.23  | Good      | E     | VS1     | 56.9  | 65    | 327   | 4.05  | 4.07  | 2.31  |
| 4 | 0.29  | Premium   | I     | VS2     | 62.4  | 58    | 334   | 4.2   | 4.23  | 2.63  |
| 5 | 0.31  | Good      | J     | SI2     | 63.3  | 58    | 335   | 4.34  | 4.35  | 2.75  |
| 6 | 0.24  | Very Good | J     | VVS2    | 62.8  | 57    | 336   | 3.94  | 3.96  | 2.48  |

## Selecting a single column of a **data.frame**:

To grab just the values from a single column, you would use the `pull` function. The output will be a vector (and not a **tibble**).

Since this is a long vector we will just show the first 6 values using the `head` function around the output of the `pull` function.

```
head(pull(diamonds,carat))
```

```
[1] 0.23 0.21 0.23 0.29 0.31 0.24
```

## Using the **pipe** (comes with **dplyr**):

That was a lot of typing and nested functions, which can be confusing. The pipe `%>%` or (you might see also this pipe `|>`) makes things such as this much more readable. It reads left side “pipes” into right side. RStudio **CMD/Ctrl + Shift + M** shortcut.

## Using the **pipe** (comes with **dplyr**):

Pipe `diamonds` into `select`, then pipe that into `pull`, and then show the head:

```
diamonds %>% pull(carat) %>% head()
```

```
[1] 0.23 0.21 0.23 0.29 0.31 0.24
```

## Selecting a single column of a **data.frame**:

The `pull` function is equivalent to using the `$` method (in base R).

Note that base R and tidyverse don't always play nice together.

```
head(pull(diamonds, carat))
```

```
[1] 0.23 0.21 0.23 0.29 0.31 0.24
```

```
head(diamonds$carat)
```

```
[1] 0.23 0.21 0.23 0.29 0.31 0.24
```

Note this does *not* return a **tibble** (or **data.frame**) but rather a vector.



## Selecting a single column of a `data.frame`:

The `select` function extracts one or more columns from a `tibble` or `data.frame` and returns a `tibble` (not a vector).

```
select(diamonds, carat)
```

```
# A tibble: 53,940 × 1
```

```
  carat
```

```
<dbl>
```

```
1  0.23
```

```
2  0.21
```

```
3  0.23
```

```
4  0.29
```

```
5  0.31
```

```
6  0.24
```

```
7  0.24
```

```
8  0.26
```

```
9  0.22
```

```
10 0.23
```

```
# i 53,930 more rows
```

## Selecting multiple columns of a `data.frame`:

The `select` command from `dplyr` is very flexible. You just need to list all columns you want to extract separated by commas. You can use this as a way to just keep the columns you want for example.

```
select(diamonds, carat, depth)
```

```
# A tibble: 53,940 × 2
```

```
  carat depth  
  <dbl> <dbl>  
1  0.23  61.5  
2  0.21  59.8  
3  0.23  56.9  
4  0.29  62.4  
5  0.31  63.3  
6  0.24  62.8  
7  0.24  62.3  
8  0.26  61.9  
9  0.22  65.1  
10 0.23  59.4
```

```
# i 53,930 more rows
```

## See the Select “helpers”

Type `tidyselect::` to see functions available.



Here are a few:

```
last_col()  
ends_with()  
starts_with()  
contains() # search for a pattern  
everything()
```

# Tidysselect helpers

For example, we can take all columns that start with a "c":

```
diamonds %>% select(starts_with("c"))
```

```
# A tibble: 53,940 × 4
  carat cut      color clarity
  <dbl> <ord>    <ord> <ord>
1  0.23 Ideal    E     SI2
2  0.21 Premium  E     SI1
3  0.23 Good     E     VS1
4  0.29 Premium  I     VS2
5  0.31 Good     J     SI2
6  0.24 Very Good J     VVS2
7  0.24 Very Good I     VVS1
8  0.26 Very Good H     SI1
9  0.22 Fair     E     VS2
10 0.23 Very Good H     VS1
# i 53,930 more rows
```

# Tidysselect helpers

Or we can take all columns that end with an "e":

```
diamonds %>% select(ends_with("e"))
```

```
# A tibble: 53,940 × 2
```

```
  table price
```

```
  <dbl> <int>
```

|    |    |     |
|----|----|-----|
| 1  | 55 | 326 |
| 2  | 61 | 326 |
| 3  | 65 | 327 |
| 4  | 58 | 334 |
| 5  | 58 | 335 |
| 6  | 57 | 336 |
| 7  | 57 | 336 |
| 8  | 55 | 337 |
| 9  | 61 | 337 |
| 10 | 61 | 338 |

```
# i 53,930 more rows
```

## Tidymodels helpers

We are going to cover “fancier” ways of matching column names (and strings more generally) in the data cleaning lecture.

## Subset rows of a `data.frame`:

The command in `dplyr` for subsetting rows is `filter`. Try `?filter`.

The easiest way to filter is by testing whether numeric observations are greater than or less than some cutoff:

```
filter(diamonds, depth > 60)
```

```
# A tibble: 48,315 × 10
```

|    | carat | cut       | color | clarity | depth | table | price | x     | y     | z     |
|----|-------|-----------|-------|---------|-------|-------|-------|-------|-------|-------|
|    | <dbl> | <ord>     | <ord> | <ord>   | <dbl> | <dbl> | <int> | <dbl> | <dbl> | <dbl> |
| 1  | 0.23  | Ideal     | E     | SI2     | 61.5  | 55    | 326   | 3.95  | 3.98  | 2.43  |
| 2  | 0.29  | Premium   | I     | VS2     | 62.4  | 58    | 334   | 4.2   | 4.23  | 2.63  |
| 3  | 0.31  | Good      | J     | SI2     | 63.3  | 58    | 335   | 4.34  | 4.35  | 2.75  |
| 4  | 0.24  | Very Good | J     | VVS2    | 62.8  | 57    | 336   | 3.94  | 3.96  | 2.48  |
| 5  | 0.24  | Very Good | I     | VVS1    | 62.3  | 57    | 336   | 3.95  | 3.98  | 2.47  |
| 6  | 0.26  | Very Good | H     | SI1     | 61.9  | 55    | 337   | 4.07  | 4.11  | 2.53  |
| 7  | 0.22  | Fair      | E     | VS2     | 65.1  | 61    | 337   | 3.87  | 3.78  | 2.49  |
| 8  | 0.3   | Good      | J     | SI1     | 64    | 55    | 339   | 4.25  | 4.28  | 2.73  |
| 9  | 0.23  | Ideal     | J     | VS1     | 62.8  | 56    | 340   | 3.93  | 3.9   | 2.46  |
| 10 | 0.22  | Premium   | F     | SI1     | 60.4  | 61    | 342   | 3.88  | 3.84  | 2.33  |

```
# i 48,305 more rows
```

Note, no subsetting is necessary. R “knows” `depth` refers to a column of `diamonds`.

## Subset using pipes:

You can also using piping here:

```
diamonds %>% filter(depth > 60)
```

```
# A tibble: 48,315 × 10
```

|    | carat | cut       | color | clarity | depth | table | price | x     | y     | z     |
|----|-------|-----------|-------|---------|-------|-------|-------|-------|-------|-------|
|    | <dbl> | <ord>     | <ord> | <ord>   | <dbl> | <dbl> | <int> | <dbl> | <dbl> | <dbl> |
| 1  | 0.23  | Ideal     | E     | SI2     | 61.5  | 55    | 326   | 3.95  | 3.98  | 2.43  |
| 2  | 0.29  | Premium   | I     | VS2     | 62.4  | 58    | 334   | 4.2   | 4.23  | 2.63  |
| 3  | 0.31  | Good      | J     | SI2     | 63.3  | 58    | 335   | 4.34  | 4.35  | 2.75  |
| 4  | 0.24  | Very Good | J     | VVS2    | 62.8  | 57    | 336   | 3.94  | 3.96  | 2.48  |
| 5  | 0.24  | Very Good | I     | VVS1    | 62.3  | 57    | 336   | 3.95  | 3.98  | 2.47  |
| 6  | 0.26  | Very Good | H     | SI1     | 61.9  | 55    | 337   | 4.07  | 4.11  | 2.53  |
| 7  | 0.22  | Fair      | E     | VS2     | 65.1  | 61    | 337   | 3.87  | 3.78  | 2.49  |
| 8  | 0.3   | Good      | J     | SI1     | 64    | 55    | 339   | 4.25  | 4.28  | 2.73  |
| 9  | 0.23  | Ideal     | J     | VS1     | 62.8  | 56    | 340   | 3.93  | 3.9   | 2.46  |
| 10 | 0.22  | Premium   | F     | SI1     | 60.4  | 61    | 342   | 3.88  | 3.84  | 2.33  |

```
# i 48,305 more rows
```



## Subset rows with multiple required conditions:

You can combine filtering on multiple columns by separating the filter arguments with & (or commas - but & is more clear)

```
diamonds %>% filter(depth > 60 & table > 60 & price > 2775)
```

```
# A tibble: 1,704 × 10
```

|    | carat<br><dbl> | cut<br><ord> | color<br><ord> | clarity<br><ord> | depth<br><dbl> | table<br><dbl> | price<br><int> | x<br><dbl> | y<br><dbl> | z<br><dbl> |
|----|----------------|--------------|----------------|------------------|----------------|----------------|----------------|------------|------------|------------|
| 1  | 0.72           | Premium      | F              | SI1              | 61.8           | 61             | 2777           | 5.82       | 5.71       | 3.56       |
| 2  | 0.72           | Very Good    | H              | VS1              | 60.6           | 63             | 2782           | 5.83       | 5.76       | 3.51       |
| 3  | 0.81           | Good         | G              | SI2              | 61             | 61             | 2789           | 5.94       | 5.99       | 3.64       |
| 4  | 0.71           | Premium      | F              | VS1              | 60.1           | 62             | 2790           | 5.77       | 5.74       | 3.46       |
| 5  | 0.71           | Premium      | G              | VS1              | 62.4           | 61             | 2803           | 5.7        | 5.65       | 3.54       |
| 6  | 0.74           | Fair         | F              | VS2              | 61.1           | 68             | 2805           | 5.82       | 5.75       | 3.53       |
| 7  | 0.7            | Good         | F              | VS1              | 62.8           | 61             | 2810           | 5.57       | 5.61       | 3.51       |
| 8  | 0.7            | Very Good    | F              | VS2              | 60.9           | 61             | 2812           | 5.66       | 5.71       | 3.46       |
| 9  | 0.71           | Good         | E              | SI1              | 62.8           | 64             | 2817           | 5.6        | 5.54       | 3.5        |
| 10 | 0.7            | Premium      | E              | VS2              | 62.4           | 61             | 2818           | 5.66       | 5.63       | 3.52       |

```
# i 1,694 more rows
```

## Subset rows for a character condition:

You can also filter character strings by a single value or category. Here we need quotes around character strings.

```
diamonds %>% filter(color == "I" &  
  clarity == "SI2" & cut == "Premium")
```

```
# A tibble: 312 × 10
```

|    | carat<br><dbl> | cut<br><ord> | color<br><ord> | clarity<br><ord> | depth<br><dbl> | table<br><dbl> | price<br><int> | x<br><dbl> | y<br><dbl> | z<br><dbl> |
|----|----------------|--------------|----------------|------------------|----------------|----------------|----------------|------------|------------|------------|
| 1  | 0.42           | Premium      | I              | SI2              | 61.5           | 59             | 552            | 4.78       | 4.84       | 2.96       |
| 2  | 1              | Premium      | I              | SI2              | 58.2           | 60             | 2795           | 6.61       | 6.55       | 3.83       |
| 3  | 0.9            | Premium      | I              | SI2              | 62.2           | 59             | 2826           | 6.11       | 6.07       | 3.79       |
| 4  | 1.05           | Premium      | I              | SI2              | 58.3           | 57             | 2911           | 6.72       | 6.67       | 3.9        |
| 5  | 0.91           | Premium      | I              | SI2              | 62             | 59             | 2913           | 6.18       | 6.23       | 3.85       |
| 6  | 0.9            | Premium      | I              | SI2              | 62.5           | 58             | 2948           | 6.15       | 6.1        | 3.83       |
| 7  | 0.9            | Premium      | I              | SI2              | 60.6           | 60             | 2948           | 6.28       | 6.23       | 3.79       |
| 8  | 1.06           | Premium      | I              | SI2              | 61.5           | 57             | 2968           | 6.57       | 6.49       | 4.02       |
| 9  | 0.91           | Premium      | I              | SI2              | 60.2           | 59             | 2981           | 6.29       | 6.24       | 3.77       |
| 10 | 0.9            | Premium      | I              | SI2              | 60.6           | 60             | 3001           | 6.23       | 6.28       | 3.79       |

```
# i 302 more rows
```

## Subset to match a vector of possibilities:

Sometimes you want to be able to filter on matching several values or categories. The `%in%` operator is useful here:

```
diamonds %>% filter(clarity %in% c("SI1", "SI2"))
```

```
# A tibble: 22,259 × 10
```

|    | carat<br><dbl> | cut<br><ord> | color<br><ord> | clarity<br><ord> | depth<br><dbl> | table<br><dbl> | price<br><int> | x<br><dbl> | y<br><dbl> | z<br><dbl> |
|----|----------------|--------------|----------------|------------------|----------------|----------------|----------------|------------|------------|------------|
| 1  | 0.23           | Ideal        | E              | SI2              | 61.5           | 55             | 326            | 3.95       | 3.98       | 2.43       |
| 2  | 0.21           | Premium      | E              | SI1              | 59.8           | 61             | 326            | 3.89       | 3.84       | 2.31       |
| 3  | 0.31           | Good         | J              | SI2              | 63.3           | 58             | 335            | 4.34       | 4.35       | 2.75       |
| 4  | 0.26           | Very Good    | H              | SI1              | 61.9           | 55             | 337            | 4.07       | 4.11       | 2.53       |
| 5  | 0.3            | Good         | J              | SI1              | 64             | 55             | 339            | 4.25       | 4.28       | 2.73       |
| 6  | 0.22           | Premium      | F              | SI1              | 60.4           | 61             | 342            | 3.88       | 3.84       | 2.33       |
| 7  | 0.31           | Ideal        | J              | SI2              | 62.2           | 54             | 344            | 4.35       | 4.37       | 2.71       |
| 8  | 0.2            | Premium      | E              | SI2              | 60.2           | 62             | 345            | 3.79       | 3.75       | 2.27       |
| 9  | 0.3            | Ideal        | I              | SI2              | 62             | 54             | 348            | 4.31       | 4.34       | 2.68       |
| 10 | 0.3            | Good         | J              | SI1              | 63.4           | 54             | 351            | 4.23       | 4.29       | 2.7        |

```
# i 22,249 more rows
```

## Subset rows using conditions for numeric and character variables:

You can mix and match filtering on numeric and categorical/character columns in the same `filter()` command:

```
diamonds %>% filter(clarity %in% c("SI1", "SI2") &  
                    cut == "Premium" & price > 3000)
```

```
# A tibble: 3,976 × 10
```

|    | carat<br><dbl> | cut<br><ord> | color<br><ord> | clarity<br><ord> | depth<br><dbl> | table<br><dbl> | price<br><int> | x<br><dbl> | y<br><dbl> | z<br><dbl> |
|----|----------------|--------------|----------------|------------------|----------------|----------------|----------------|------------|------------|------------|
| 1  | 0.9            | Premium      | I              | SI2              | 60.6           | 60             | 3001           | 6.23       | 6.28       | 3.79       |
| 2  | 0.81           | Premium      | F              | SI1              | 61.9           | 58             | 3004           | 5.99       | 5.96       | 3.7        |
| 3  | 0.92           | Premium      | D              | SI2              | 60.2           | 61             | 3004           | 6.32       | 6.27       | 3.79       |
| 4  | 0.9            | Premium      | D              | SI1              | 62.2           | 60             | 3013           | 6.08       | 6.05       | 3.77       |
| 5  | 0.96           | Premium      | E              | SI2              | 62.8           | 60             | 3016           | 6.3        | 6.24       | 3.94       |
| 6  | 0.93           | Premium      | G              | SI2              | 61.4           | 56             | 3019           | 6.27       | 6.23       | 3.84       |
| 7  | 0.78           | Premium      | D              | SI1              | 60.4           | 57             | 3019           | 6.02       | 5.97       | 3.62       |
| 8  | 0.75           | Premium      | E              | SI1              | 61.7           | 60             | 3024           | 5.84       | 5.8        | 3.59       |
| 9  | 0.75           | Premium      | D              | SI1              | 59.2           | 58             | 3024           | 5.96       | 5.93       | 3.52       |
| 10 | 1.02           | Premium      | G              | SI2              | 61.7           | 58             | 3027           | 6.46       | 6.41       | 3.97       |

```
# i 3,966 more rows
```

## Note about quotes and numbers

R will interpret quotes around numbers as the characters themselves and not their numeric meaning. Thus it's generally best to avoid quotes around numeric unless it is not being treated as a numeric value - for example levels or grades.

```
diamonds %>% filter(price > 3001) #This works  
diamonds %>% filter(price > "3001") # This does not  
  
diamonds %>% filter(price == 3001) # This works  
diamonds %>% filter(price == "3001") # this works
```

## Logic for subsetting rows of a **data.frame**:

Other useful logical tests:

**&** : AND

**|** : OR

**<=** : less than or equals

**>=** : greater than or equals

**!=** : not equals

## Subset rows with or logic:

The OR operator (|) is more permissive than the AND operator:

```
diamonds %>% filter(depth > 60 | table > 60 | price > 2775)
```

```
# A tibble: 52,198 × 10
```

|    | carat | cut       | color | clarity | depth | table | price | x     | y     | z     |
|----|-------|-----------|-------|---------|-------|-------|-------|-------|-------|-------|
|    | <dbl> | <ord>     | <ord> | <ord>   | <dbl> | <dbl> | <int> | <dbl> | <dbl> | <dbl> |
| 1  | 0.23  | Ideal     | E     | SI2     | 61.5  | 55    | 326   | 3.95  | 3.98  | 2.43  |
| 2  | 0.21  | Premium   | E     | SI1     | 59.8  | 61    | 326   | 3.89  | 3.84  | 2.31  |
| 3  | 0.23  | Good      | E     | VS1     | 56.9  | 65    | 327   | 4.05  | 4.07  | 2.31  |
| 4  | 0.29  | Premium   | I     | VS2     | 62.4  | 58    | 334   | 4.2   | 4.23  | 2.63  |
| 5  | 0.31  | Good      | J     | SI2     | 63.3  | 58    | 335   | 4.34  | 4.35  | 2.75  |
| 6  | 0.24  | Very Good | J     | VVS2    | 62.8  | 57    | 336   | 3.94  | 3.96  | 2.48  |
| 7  | 0.24  | Very Good | I     | VVS1    | 62.3  | 57    | 336   | 3.95  | 3.98  | 2.47  |
| 8  | 0.26  | Very Good | H     | SI1     | 61.9  | 55    | 337   | 4.07  | 4.11  | 2.53  |
| 9  | 0.22  | Fair      | E     | VS2     | 65.1  | 61    | 337   | 3.87  | 3.78  | 2.49  |
| 10 | 0.23  | Very Good | H     | VS1     | 59.4  | 61    | 338   | 4     | 4.05  | 2.39  |

```
# i 52,188 more rows
```

## OR is like %in%:

The OR operator (|) can be a substitute for %in% (although it might take more typing):

```
diamonds %>% filter(clarity == "SI1" | clarity == "SI2") %>% head(2)
```

```
# A tibble: 2 × 10
  carat cut      color clarity depth table price      x      y      z
  <dbl> <ord>    <ord> <ord>    <dbl> <dbl> <int> <dbl> <dbl> <dbl>
1  0.23 Ideal    E      SI2     61.5    55   326   3.95   3.98   2.43
2  0.21 Premium E      SI1     59.8    61   326   3.89   3.84   2.31
```

```
diamonds %>% filter(clarity %in% c("SI1", "SI2")) %>% head(2)
```

```
# A tibble: 2 × 10
  carat cut      color clarity depth table price      x      y      z
  <dbl> <ord>    <ord> <ord>    <dbl> <dbl> <int> <dbl> <dbl> <dbl>
1  0.23 Ideal    E      SI2     61.5    55   326   3.95   3.98   2.43
2  0.21 Premium E      SI1     59.8    61   326   3.89   3.84   2.31
```



## Combining **filter** and **select**:

You can combine **filter** and **select** to subset the rows and columns, respectively, of a **data.frame**:

```
diamonds %>%  
  filter(clarity == "SI2") %>%  
  select(starts_with("c"))
```

```
# A tibble: 9,194 × 4  
  carat cut      color clarity  
  <dbl> <ord>    <ord> <ord>  
1  0.23 Ideal    E      SI2  
2  0.31 Good     J      SI2  
3  0.31 Ideal    J      SI2  
4  0.2  Premium  E      SI2  
5  0.3  Ideal    I      SI2  
6  0.3  Good     I      SI2  
7  0.33 Ideal    I      SI2  
8  0.33 Ideal    I      SI2  
9  0.32 Good     H      SI2  
10 0.32 Very Good H      SI2  
# i 9,184 more rows
```

## Combining **filter** and **select**- Order matters!:

The order of these functions matters though, since you can remove columns that you might want to filter on.

```
diamonds %>%  
  select(starts_with("c")) %>%  
  filter(table > 60))
```

This will result in an error because the table column is now gone after the `select()` function!

Fancier filtering

## Combining tidysselect helpers with regular selection

```
head(diamonds, 2)
```

```
# A tibble: 2 × 10
  carat cut      color clarity depth table price      x      y      z
  <dbl> <ord>    <ord> <ord>    <dbl> <dbl> <int> <dbl> <dbl> <dbl>
1  0.23 Ideal    E      SI2     61.5   55   326  3.95  3.98  2.43
2  0.21 Premium E      SI1     59.8   61   326  3.89  3.84  2.31
```

```
diamonds %>% select(price, starts_with("c"))
```

```
# A tibble: 53,940 × 5
  price carat cut      color clarity
  <int> <dbl> <ord>    <ord> <ord>
1    326  0.23 Ideal    E      SI2
2    326  0.21 Premium  E      SI1
3    327  0.23 Good     E      VS1
4    334  0.29 Premium  I      VS2
5    335  0.31 Good     J      SI2
6    336  0.24 Very Good J      VVS2
7    336  0.24 Very Good I      VVS1
8    337  0.26 Very Good H      SI1
9    337  0.22 Fair     E      VS2
10   338  0.23 Very Good H      VS1
# i 53,930 more rows
```

# Multiple tidyselect functions

Follows OR logic.

```
diamonds %>% select(starts_with("c"), ends_with("e"))
```

```
# A tibble: 53,940 × 6
```

|    | carat | cut       | color | clarity | table | price |
|----|-------|-----------|-------|---------|-------|-------|
|    | <dbl> | <ord>     | <ord> | <ord>   | <dbl> | <int> |
| 1  | 0.23  | Ideal     | E     | SI2     | 55    | 326   |
| 2  | 0.21  | Premium   | E     | SI1     | 61    | 326   |
| 3  | 0.23  | Good      | E     | VS1     | 65    | 327   |
| 4  | 0.29  | Premium   | I     | VS2     | 58    | 334   |
| 5  | 0.31  | Good      | J     | SI2     | 58    | 335   |
| 6  | 0.24  | Very Good | J     | VVS2    | 57    | 336   |
| 7  | 0.24  | Very Good | I     | VVS1    | 57    | 336   |
| 8  | 0.26  | Very Good | H     | SI1     | 55    | 337   |
| 9  | 0.22  | Fair      | E     | VS2     | 61    | 337   |
| 10 | 0.23  | Very Good | H     | VS1     | 61    | 338   |

```
# i 53,930 more rows
```

## Multiple patterns with tidyselect

Need to combine the patterns with the `c()` function.

```
diamonds %>% select(starts_with(c("c", "p")))
```

```
# A tibble: 53,940 × 5
  carat cut      color clarity price
  <dbl> <ord>    <ord> <ord>    <int>
1  0.23 Ideal    E      SI2      326
2  0.21 Premium  E      SI1      326
3  0.23 Good     E      VS1      327
4  0.29 Premium  I      VS2      334
5  0.31 Good     J      SI2      335
6  0.24 Very Good J      VVS2     336
7  0.24 Very Good I      VVS1     336
8  0.26 Very Good H      SI1      337
9  0.22 Fair     E      VS2      337
10 0.23 Very Good H      VS1      338
# i 53,930 more rows
```

## Common error for filter or select

If you try to filter or select for a column that does not exist it will not work:

- misspelled column name
- column that was already removed

**Always good to check each step!**

Did the filter work the way you expected? Did the dimensions change?

Don't use quotes on column names - nothing will get filtered!



Source: <https://media.giphy.com/media/5b5OU7aUekfdSAER5I/giphy.gif>



# Summary

- `pull()` can help us see a vector version of our variables - we can “pull” out the data from a dataframe
- The pipe (`%>%`) can help us to do sequential steps
- `select()` makes a smaller table of just selected variables
- tidyselect functions can help us select specific columns: `contains()`, `ends_with()`, `starts_with()`
- `filter` can remove rows based on conditions
- `==` is needed to filter for rows that are “exactly equal” to a value
- `!=` does the opposite

## Summary Continued

- `%in%` enables us to do multiple `==` conditions such as `%in% c(1,2,3)`
- `|` is for or logic and `&` is for and logic when combining filter conditions together
- Always check that you filtered for what you think you did!

# Lab

[Link to Lab](#)