

Data Cleaning Part 2

Data Wrangling in R

Data Cleaning Part 2

Example of Cleaning: more complicated

For example, let's say we have a variable about treatment or control conditions coded as treatment, T, treat, Treat, C, Cont, cont, cOnt, Control, and control. Using Excel to find all of these would be a matter of filtering and changing all by hand or using if statements.

Sometimes though, it's not so simple. That's where functions that find patterns come to be very useful.

Some toy data for examples

```
set.seed(4) # random sample below – make sure same every time
status <- sample(c("treatment", "T", "treat",
                  "Traet", "Treat", "C", "Cont",
                  "cont", "cOnt", "Control", "control"),
                1000, replace = TRUE)
data_case = tibble(status)
```

Take a look at the data

```
count(data_case, status)
```

```
# A tibble: 11 × 2
  status      n
  <chr>    <int>
1 C         81
2 Cont      90
3 Control   91
4 T         91
5 Traet    105
6 Treat    100
7 cOnt      79
8 cont      83
9 control   98
10 treat     86
11 treatment 96
```

Example of Cleaning: more complicated

In R, you could use `case_when()`:

```
#case_when way:
data_case <- data_case %>% mutate(status =
  case_when(status
    %in% c("C", "cont", "cOnt", "Cont", "control", "Control")
          ~ "Control",
    .default = status))
count(data_case, status)
```

```
# A tibble: 6 × 2
  status      n
  <chr>    <int>
1 Control    522
2 T          91
3 Traet     105
4 Treat     100
5 treat      86
6 treatment  96
```

Oh dear! This only fixes some values! It is difficult to notice values like "Traet".

String functions

The **stringr** package

Like `dplyr`, the `stringr` package:

- Makes some things more intuitive
- Is different than base R
- Is used on forums for answers
- Has a standard format for most functions: `str_`
 - the first argument is a string like first argument is a `data.frame` in `dplyr`

Useful String Functions

Useful String functions from base R and `stringr`

- `toupper()`, `tolower()` - uppercase or lowercase your data
- `str_sentence()` - uppercase just the first character (in the `stringr` package)
- `paste()` - paste strings together with a space
- `paste0` - paste strings together with no space as default
- `str_trim()` (in the `stringr` package) or `trimws` in base
 - will trim whitespace
- `nchar` - get the number of characters in a string

recoding with `str_to_sentence()`

```
data_case <- data_case %>%  
  mutate(status = str_to_sentence(status))  
count(data_case, status)
```

```
# A tibble: 5 × 2  
  status      n  
  <chr>    <int>  
1 Control    522  
2 T          91  
3 Traet     105  
4 Treat     186  
5 Treatment   96
```

recoding with `str_to_sentence()`

```
#case_when way:
data_case <-data_case %>%
  mutate(status = str_to_sentence(status)) %>%
  mutate(status =
    case_when(status %in%
      c("Treatment", "T", "Treat", "Traet", "Treat")
        ~ "Treatment",
      .default = status))
count(data_case, status)
```

```
# A tibble: 2 × 2
  status      n
  <chr>    <int>
1 Control    522
2 Treatment  478
```

OK, now we are getting somewhere!

Reading in again

Now we have a chance to keep but clean these values!

```
ufo <-read_csv(  
  "https://sisbid.github.io/Data-Wrangling/data/ufo/ufo_data_complete.csv",  
  col_types = cols(`duration (seconds)` = "c"))
```

Warning: One or more parsing issues, call `problems()` on your data frame for
dat <- vroom(...)
problems(dat)

Clean names with the `clean_names()` function from the **janitor** package

```
colnames(ufo)
```

```
[1] "datetime"      "city"          "state"
[4] "country"       "shape"         "duration (seconds)"
[7] "duration (hours/min)" "comments"      "date posted"
[10] "latitude"      "longitude"
```

```
ufo_clean <- clean_names(ufo)
colnames(ufo_clean)
```

```
[1] "datetime"      "city"          "state"
[4] "country"       "shape"         "duration_seconds"
[7] "duration_hours_min" "comments"      "date_posted"
[10] "latitude"      "longitude"
```

str_detect and filter

Now let's fix our ufo data and remove those pesky backticks in the `duration_seconds` variable. First let's find them with `str_detect`.

```
ufo_clean %>%  
  filter(str_detect(  
    string = duration_seconds,  
    pattern = "`"))
```

```
# A tibble: 3 × 11  
  datetime      city state country shape duration_seconds duration_hours_min  
  <chr>         <chr> <chr> <chr>   <chr> <chr>          <chr>  
1 2/2/2000 19:33 bouse az    us    <NA> 2`          each a few seconds  
2 4/10/2005 22:52 santa... ca    us    <NA> 8`          eight seconds  
3 7/21/2006 13:00 ibagu... <NA> <NA>   circ... 0.5`        1/2 segundo  
# i 4 more variables: comments <chr>, date_posted <chr>, latitude <chr>,  
# longitude <chr>
```

str_remove

Now, let's remove the backticks.

```
ufo_clean <- ufo_clean %>%  
  mutate(duration_seconds =  
    str_remove(string = duration_seconds,  
               pattern = "`"))
```

And let's check it:

```
ufo_clean %>%  
  filter(str_detect(  
    string = duration_seconds,  
    pattern = "`"))
```

```
# A tibble: 0 × 11  
# i 11 variables: datetime <chr>, city <chr>, state <chr>, country <chr>,  
#   shape <chr>, duration_seconds <chr>, duration_hours_min <chr>,  
#   comments <chr>, date_posted <chr>, latitude <chr>, longitude <chr>
```

Lets also mutate to be numeric again

There are many class functions to change from one class to another start with "as."

```
ufo_clean <- ufo_clean %>%  
  mutate(duration_seconds = as.numeric(duration_seconds))  
ufo_clean
```

```
# A tibble: 88,875 × 11  
  datetime      city state country shape duration_seconds duration_hours_min  
  <chr>         <chr> <chr> <chr>   <chr>          <dbl> <chr>  
1 10/10/1949 20:... san ... tx    us    cyli...      2700 45 minutes  
2 10/10/1949 21:... lack... tx    <NA>    light      7200 1-2 hrs  
3 10/10/1955 17:... ches... <NA>   gb     circ...       20 20 seconds  
4 10/10/1956 21:... edna  tx    us     circ...       20 1/2 hour  
5 10/10/1960 20:... kane... hi    us     light      900 15 minutes  
6 10/10/1961 19:... bris... tn    us     sphe...      300 5 minutes  
7 10/10/1965 21:... pena... <NA>   gb     circ...      180 about 3 mins  
8 10/10/1965 23:... norw... ct    us     disk      1200 20 minutes  
9 10/10/1966 20:... pell... al    us     disk       180 3 minutes  
10 10/10/1966 21:... live... fl    us     disk       120 several minutes  
# i 88,865 more rows  
# i 4 more variables: comments <chr>, date_posted <chr>, latitude <chr>,  
# longitude <chr>
```


Substringing

stringr

- `str_sub(x, start, end)` - substrings from position start to position end

Substringing

Examples:

```
str_sub("I like friesian horses", 8,12)
```

```
[1] "fries"
```

```
#123456789101112
```

```
#I like fries
```

```
str_sub(c("Site A", "Site B", "Site C"), 6,6)
```

```
[1] "A" "B" "C"
```

Splitting/Find/Replace and Regular Expressions

- R can do much more than find exact matches for a whole string
- Like Perl and other languages, it can use regular expressions.
- What are regular expressions?
 - Ways to search for specific strings
 - Can be very complicated or simple
 - Highly Useful - think “Find” on steroids

A bit on Regular Expressions

- <https://rstudio.github.io/cheatsheets/strings.pdf>
- They can use to match a large number of strings in one statement
- `.` matches any single character
- `*` means repeat as many (even if 0) more times the last character
- `?` makes a pattern optional (i.e. it matches 0 or 1 times)
- `^` matches start of vector `^a` - starts with "a"
- `$` matches end of vector `b$` - ends with "b"

'Find' functions: **stringr**

`str_detect`, `str_subset`, `str_replace`, and `str_replace_all` search for matches to argument pattern within each element of a character vector: they differ in the format of and amount of detail in the results.

- `str_detect` - returns TRUE if pattern is found
- `str_subset` - returns only the strings where the pattern were detected
- `str_extract` - returns only the pattern that was detected
- `str_replace` - replaces pattern with replacement the first time
- `str_replace_all` - replaces pattern with replacement as many times matched

'Find' functions: Finding Indices

These are the indices where the pattern match occurs:

```
ufo_clean %>%  
  filter(str_detect(comments, "two aliens")) %>%  
  head()  
  
# A tibble: 2 × 11  
  datetime          city state country shape duration_seconds duration_hours_min  
  <chr>            <chr> <chr> <chr>   <chr>          <dbl> <chr>  
1 10/14/2006 02:00 yuma  va    us    form...      300 5 minutes  
2 7/1/2007 23:00  nort... ct    <NA>   unkn...      60 1 minute  
# i 4 more variables: comments <chr>, date_posted <chr>, latitude <chr>,  
# longitude <chr>
```

To Take a look at comments... need to select it first

```
ufo_clean %>%  
  filter(str_detect(comments, "two aliens")) %>%  
  select(comments)
```

```
# A tibble: 2 × 1
```

```
  comments
```

```
  <chr>
```

```
1 ((HOAX??)) two aliens appeared from a bright light to peacefully investigate...
```

```
2 Witnessed two aliens walking along baseball field fence.
```

'Find' functions: `str_subset()` is easier

`str_subset()` gives the values that match the pattern: (or if we used `negate = TRUE` it would find the opposite)

```
ufo_clean %>% pull(comments) %>%  
  str_subset( "two aliens")
```

```
[1] "((HOAX??))  two aliens appeared from a bright light to peacefully investigate the surroundings in the woods"  
[2] "Witnessed two aliens walking along baseball field fence."
```


Showing difference in `str_extract`

`str_extract` extracts just the matched string

```
ufo_clean %>%  
  mutate(aliens = str_extract(comments, "two aliens")) %>%  
  count(aliens)
```

```
# A tibble: 2 × 2  
  aliens      n  
  <chr>    <int>  
1 two aliens      2  
2 <NA>      88873
```

- Look for any comment that starts with “aliens”

```
ufo_clean %>% pull(comments) %>% str_subset( "^aliens")
```

```
[1] "aliens speak german???" "aliens exist"          "aliens in srilanka"
```

Using Regular Expressions

That contains space then ship maybe with stuff in between

```
ufo_clean %>% pull(comments) %>%  
  str_subset("space.?ship") %>% head(4) # gets "spaceship" or "space ship" or...
```

```
[1] "I saw the cylinder shaped looked like a spaceship hovering above the east side of the Air Force base. Saw it for  
[2] "description of a spaceship spotted over Birmingham Alabama in 1967."  
[3] "A space ship was descending to the ground"  
[4] "On Monday october 3rd 2005 I spotted two spaceships in the sky. The first spotted ship was what seemed to be a
```

```
ufo_clean %>% pull(comments) %>%  
  str_subset("space.ship") %>% head(4) # no "spaceship" must have character in between
```

```
[1] "A space ship was descending to the ground"  
[2] "I saw a Silver space ship rising into the early morning sky over Houston Texas."  
[3] "Saw a space ship hanging over the southern (Manzano) portion of the Sandia Mountains on evening. It was bright  
[4] "saw space ship for 5 min Got scared crapless"
```

time information

```
pull(ufo_clean, duration_hours_min) %>% head(n = 20)
```

[1]	"45 minutes"	"1-2 hrs"	"20 seconds"	"1/2 hour"
[5]	"15 minutes"	"5 minutes"	"about 3 mins"	"20 minutes"
[9]	"3 minutes"	"several minutes"	"5 min."	"3 minutes"
[13]	"30 min."	"3 minutes"	"30 seconds"	"20minutes"
[17]	"2 minutes"	"20-30 min"	"20 sec."	"2 min"

str_replace()

Let's say we wanted to make the time information more consistent. Using `case_when()` could be very tedious and error-prone!

We can use `str_replace()` to do so.

```
ufo_clean %>% mutate(duration_hours_min =  
  str_replace(string = duration_hours_min,  
    pattern = "minutes",  
    replacement = "mins")) %>%  
  pull(duration_hours_min) %>%  
  head(8)
```

```
[1] "45 mins"      "1-2 hrs"      "20 seconds"   "1/2 hour"     "15 mins"  
[6] "5 mins"       "about 3 mins" "20 mins"
```

Separating columns

Better yet, you might notice that this data isn't tidy- there are more than two entries for each value - amount of time and unit. We could separate this using `separate()` from the `tidyr` package.

```
ufo_clean %>% separate(duration_hours_min,  
                        into = c("duration_amount", "duration_unit"),  
                        sep = " ") %>%  
  select(duration_amount, duration_unit) %>% head()
```

```
# A tibble: 6 × 2  
  duration_amount duration_unit  
    <chr>         <chr>  
1 45            minutes  
2 1-2           hrs  
3 20            seconds  
4 1/2           hour  
5 15            minutes  
6 5             minutes
```

As you can see there is still plenty of cleaning to do!

more seperating

```
ufo_clean <- ufo_clean %>% separate(datetime,  
                                     into = c("date", "time"),  
                                     sep = " ")  
ufo_clean %>% select(date, time) %>% head()
```

```
# A tibble: 6 × 2  
  date      time  
  <chr>    <chr>  
1 10/10/1949 20:30  
2 10/10/1949 21:00  
3 10/10/1955 17:00  
4 10/10/1956 21:00  
5 10/10/1960 20:00  
6 10/10/1961 19:00
```

Dates and times

The `lubridate` package is amazing for dates. Most important functions are those that look like `ymd` or `mdy` etc. They specify how a date should be interpreted.

```
library(lubridate)#need to load this one!
```

```
ufo_clean <- ufo_clean %>% mutate(date = mdy(date))  
head(ufo_clean)
```

```
# A tibble: 6 × 12  
  date      time city state country shape duration_seconds duration_hours_min  
  <date>    <chr> <chr> <chr> <chr>  <chr>          <dbl> <chr>  
1 1949-10-10 20:30 san ... tx    us    cyli...      2700 45 minutes  
2 1949-10-10 21:00 lack... tx    <NA>    light      7200 1-2 hrs  
3 1955-10-10 17:00 ches... <NA>    gb     circ...      20 20 seconds  
4 1956-10-10 21:00 edna  tx    us     circ...      20 1/2 hour  
5 1960-10-10 20:00 kane... hi    us     light      900 15 minutes  
6 1961-10-10 19:00 bris... tn    us     sphe...      300 5 minutes  
# i 4 more variables: comments <chr>, date_posted <chr>, latitude <chr>,  
# longitude <chr>
```

str_*functions

```
str_detect(string = c("abcdd", "two"), pattern = "dd")
```

```
[1] TRUE FALSE
```

```
str_subset(string = c("abcdd", "two"), pattern = "dd")
```

```
[1] "abcdd"
```

```
str_extract(string = c("abcdd", "two"), pattern = "dd")
```

```
[1] "dd" NA
```

```
str_sub(string = c("abcdd", "two"), start = 1, end = 3)
```

```
[1] "abc" "two"
```


Summary

- `stringr` package has lots of helpful functions that work on vectors or variables in a data frame
- `str_detect` helps find patterns
- `str_detect` and `filter` can help you filter data based on patterns within value
- `str_extract` helps extract a pattern
- `str_sub` extracts pieces of strings based on the position of the the characters
- `str_subset` gives the values that match a pattern
- `separate` can separate columns into two
- `^` indicates the start of a string
- `$` indicates the end of a string
- the `lubridate` package is useful for dates and times

Lab

https://sisbid.github.io/Data-Wrangling/09_Data_Cleaning/lab/data-cleaning-lab-part2.Rmd