

Data Cleaning

Data Wrangling in R

Data Cleaning

In general, data cleaning is a process of investigating your data for inaccuracies, or recoding it in a way that makes it more manageable.

MOST IMPORTANT RULE - LOOK AT YOUR DATA!

Read in the UFO dataset

Read in data or download from: http://sisbid.github.io/Data-Wrangling/data/ufo/ufo_data_complete.csv.gz

```
ufo <- read_delim("../data/ufo/ufo_data_complete.txt")
```

New names:

- `` -> `...12`

Warning: One or more parsing issues, see `problems()` for details

Rows: 88875 Columns: 12

— Column specification —

Delimiter: "\t"

chr (9): datetime, city, state, country, shape, duration (hours/min), comments

dbl (3): duration (seconds), longitude, ...12

i Use `spec()` to retrieve the full column specification for this data.

i Specify the column types or set `show_col_types = FALSE` to quiet this message.

The “problems”

You saw warning messages when reading in this dataset. We can see these with the `problems()` function from `readr`.

```
p <-problems(ufo)
p
```

```
# A tibble: 3 × 5
   row   col expected actual file
<int> <int> <chr>      <chr> <chr>
1 30894     6 a double 2`  /Users/avahoffman/Dropbox/JHSPH/Data-Wrangling_S
2 39616     6 a double 8`  /Users/avahoffman/Dropbox/JHSPH/Data-Wrangling_S
3 65125     6 a double 0.5` /Users/avahoffman/Dropbox/JHSPH/Data-Wrangling_S
```

The “problems”

These all became NA values.

```
ufo[(p$row-1),] %>% glimpse()
```

```
Rows: 3
Columns: 12
$ datetime      <chr> "2/2/2000 19:33", "4/10/2005 22:52", "7/21/2006
$ city          <chr> "bouse", "santa cruz", "ibague (colombia)"
$ state         <chr> "az", "ca", NA
$ country       <chr> "us", "us", NA
$ shape         <chr> NA, NA, "circle"
$ `duration (seconds)` <dbl> NA, NA, NA
$ `duration (hours/min)` <chr> "each a few seconds", "eight seconds", "1/2 seg
$ comments      <chr> "Driving through Plomosa Pass towards Bouse Loc
$ `date posted`  <chr> "2/16/2000", "4/16/2005", "10/30/2006"
$ latitude      <chr> "33.9325", "36.9741667", "4.440663"
$ longitude     <dbl> -114.00500, -122.02972, -75.24414
$ ...12         <dbl> NA, NA, NA
```

The “problems”

4/10/2005 22:52 santa cruz ca us
38453 el zumbador (el cobre) (venezuela)

8` eight seconds 2 red lights moving together and apart with a hazy cloaked craft inbetween. I
fireball 1200 20 minutes ORGANIZACI0acute;N RESCATE HUMBOLDT / SAR / VENEZUELA

Reading in again

Now we have a chance to keep but clean these values!

```
ufo <- read_delim("../data/ufo/ufo_data_complete.txt",  
                  col_types = cols("duration (seconds)" = col_character()))
```

New names:

- `` -> `...12`

```
ufo[(p$row-1),] %>% glimpse()
```

```
Rows: 3  
Columns: 12  
$ datetime      <chr> "2/2/2000 19:33", "4/10/2005 22:52", "7/21/2006  
$ city          <chr> "bouse", "santa cruz", "ibague (colombia)"  
$ state         <chr> "az", "ca", NA  
$ country       <chr> "us", "us", NA  
$ shape         <chr> NA, NA, "circle"  
$ `duration (seconds)` <chr> "2`", "8`", "0.5`"  
$ `duration (hours/min)` <chr> "each a few seconds", "eight seconds", "1/2 sec  
$ comments      <chr> "Driving through Plomosa Pass towards Bouse Loc  
$ `date posted`  <chr> "2/16/2000", "4/16/2005", "10/30/2006"  
$ latitude      <chr> "33.9325", "36.9741667", "4.440663"  
$ longitude     <dbl> -114.00500, -122.02972, -75.24414  
$ ...12        <dbl> NA, NA, NA
```

Clean names with the `clean_names()` function from the `janitor` package

```
colnames(ufo)
```

```
[1] "datetime"          "city"              "state"             "country"
[6] "duration (seconds)" "duration (hours/min)" "comments"          "date_posted"
[11] "longitude"         "...12"
```

```
ufo = clean_names(ufo)
colnames(ufo)
```

```
[1] "datetime"          "city"              "state"             "country"
[6] "duration_seconds"  "duration_hours_min" "comments"          "date_posted"
[11] "longitude"         "x12"
```


Recoding Variables

Example of Cleaning: more complicated

For example, let's say gender was coded as Male, M, m, Female, F, f. Using Excel to find all of these would be a matter of filtering and changing all by hand or using if statements.

Sometimes though, it's not so simple. That's where functions that find patterns come to be very useful.

```
table(gender)
```

gender	F	FeMAle	FEMALE	Fm	M	Ma	mAle	Male	MaLe	MALE	Man
	80	88	76	87	99	76	84	83	79	93	84

Example of Cleaning: more complicated

In R, you could use `case_when()`:

```
#case_when way:
data_gen <- data_gen %>% mutate(gender =
  case_when(gender %in% c("Male", "M", "m", "Man")
    ~ "Male",
    TRUE ~ gender))

head(data_gen)
```

```
# A tibble: 6 × 1
  gender
  <chr>
1 F
2 Fm
3 MaLe
4 MaLe
5 FeMAle
6 FEMALE
```

Oh dear! This only fixes some values! It is difficult to notice values like "MaLe".

String functions

The **stringr** package

Like `dplyr`, the `stringr` package:

- Makes some things more intuitive
- Is different than base R
- Is used on forums for answers
- Has a standard format for most functions: `str_`
 - the first argument is a string like first argument is a `data.frame` in `dplyr`

Useful String Functions

Useful String functions from base R and `stringr`

- `toupper()`, `tolower()` - uppercase or lowercase your data
- `str_sentence()` - uppercase just the first character (in the `stringr` package)
- `paste()` - paste strings together with a space
- `paste0` - paste strings together with no space as default
- `str_trim()` (in the `stringr` package) or `trimws` in base
 - will trim whitespace
- `nchar` - get the number of characters in a string

recoding with `str_to_sentence()`

```
#case_when way:
data_gen <-data_gen %>%
  mutate(gender = str_to_sentence(gender)) %>%
  mutate(gender =
    case_when(gender %in% c("Male", "M", "m", "Man")
              ~ "Male",
              TRUE ~ gender) )

head(data_gen)
```

```
# A tibble: 6 × 1
  gender
  <chr>
1 F
2 Fm
3 Male
4 Male
5 Female
6 Female
```

OK, now we are getting somewhere!

str_remove

Now let's fix our ufo data and remove those pesky backticks in the `duration_seconds` variable.

```
ufo <- ufo %>% mutate(duration_seconds = str_remove(string = duration_seconds,  
                                                    pattern = "`"))  
ufo <- ufo %>% mutate(duration_seconds = as.numeric(duration_seconds))  
ufo[(p$row-1),]
```

```
# A tibble: 3 × 12  
  datetime      city state country shape duration_seconds duration_hours_... co  
  <chr>         <chr> <chr> <chr>   <chr>          <dbl> <chr>              <chr>  
1 2/2/2000 19:... bouse az      us      <NA>          2 each a few seco... Dr  
2 4/10/2005 22... sant... ca      us      <NA>          8 eight seconds     2  
3 7/21/2006 13... ibag... <NA>    <NA>    circ...      0.5 1/2 segundo      Vi
```


Paste can add things back to variables

```
head(Orange)
```

	Tree	age	circumference
1	1	118	30
2	1	484	58
3	1	664	87
4	1	1004	115
5	1	1231	120
6	1	1372	142

```
Orange %>% mutate(Tree = paste(Tree, "Tree", sep = "_"))
```

	Tree	age	circumference
1	1_Tree	118	30
2	1_Tree	484	58
3	1_Tree	664	87
4	1_Tree	1004	115
5	1_Tree	1231	120
6	1_Tree	1372	142
7	1_Tree	1582	145
8	2_Tree	118	33
9	2_Tree	484	69
10	2_Tree	664	111
11	2_Tree	1004	156
12	2_Tree	1231	172
13	2_Tree	1372	203
14	2_Tree	1582	203
15	3_Tree	118	30

Paste0 doesn't need a separator

```
head(Orange)
```

	Tree	age	circumference
1	1	118	30
2	1	484	58
3	1	664	87
4	1	1004	115
5	1	1231	120
6	1	1372	142

```
Orange %>% mutate(Tree = paste0(Tree, "Tree"))
```

	Tree	age	circumference
1	1Tree	118	30
2	1Tree	484	58
3	1Tree	664	87
4	1Tree	1004	115
5	1Tree	1231	120
6	1Tree	1372	142
7	1Tree	1582	145
8	2Tree	118	33
9	2Tree	484	69
10	2Tree	664	111
11	2Tree	1004	156
12	2Tree	1231	172
13	2Tree	1372	203
14	2Tree	1582	203
15	3Tree	118	30

Substringing

stringr

- `str_sub(x, start, end)` - substrings from position start to position end
- `str_split(string, pattern)` - splits strings up - returns list! [we'll revisit in "Functional Programming"]

Substringing

Examples:

```
str_sub("I like friesian horses", 8,12)
```

```
[1] "fries"
```

```
#123456789101112
```

```
#I like fries
```

```
str_sub(c("Site A", "Site B", "Site C"), 6,6)
```

```
[1] "A" "B" "C"
```

Splitting/Find/Replace and Regular Expressions

- R can do much more than find exact matches for a whole string
- Like Perl and other languages, it can use regular expressions.
- What are regular expressions?
 - Ways to search for specific strings
 - Can be very complicated or simple
 - Highly Useful - think “Find” on steroids

A bit on Regular Expressions

- <http://www.regular-expressions.info/reference.html>
- They can use to match a large number of strings in one statement
- `.` matches any single character
- `*` means repeat as many (even if 0) more times the last character
- `?` makes the last thing optional
- `^` matches start of vector `^a` - starts with "a"
- `$` matches end of vector `b$` - ends with "b"

'Find' functions: **stringr**

`str_detect`, `str_subset`, `str_replace`, and `str_replace_all` search for matches to argument `pattern` within each element of a character vector: they differ in the format of and amount of detail in the results.

- `str_detect` - returns `TRUE` if `pattern` is found
- `str_subset` - returns only the strings which `pattern` were detected
 - convenient wrapper around `x[str_detect(x, pattern)]`
- `str_extract` - returns only strings which `pattern` were detected, but ONLY the `pattern`
- `str_replace` - replaces `pattern` with `replacement` the first time
- `str_replace_all` - replaces `pattern` with `replacement` as many times matched

'Find' functions: Finding Indices

These are the indices where the pattern match occurs:

```
str_detect(ufo$comments, "two aliens") %>% head()
```

```
[1] FALSE FALSE FALSE FALSE FALSE FALSE
```

```
str_detect(ufo$comments, "two aliens") %>% table()
```

```
.  
FALSE  TRUE  
88747   2
```

```
which(str_detect(ufo$comments, "two aliens"))
```

```
[1] 1730 61724
```


'Find' functions: Finding Logicals

`filter()` using `str_detect()` gives a tibble:

```
filter(ufo, str_detect(comments, "two aliens"))
```

```
# A tibble: 2 × 12
  datetime      city state country shape duration_seconds duration_hours_... comments date_posted latitude longitude
  <chr>         <chr> <chr> <chr>   <chr>          <dbl> <chr>          <chr>    <chr>      <chr>      <dbl>
1 10/14/2006 2... yuma va      us      form...      300 5 minutes      ((HOAX?... 4/27/2007    36.615      -82.6
2 7/1/2007 23:... nort... ct      <NA>      unkn...      60 1 minute      Witness... 10/19/2011   41.9856... -71.9
```

```
filter(ufo, str_detect(comments, "two aliens")) %>% select(comments)
```

```
# A tibble: 2 × 1
  comments
  <chr>
1 ((HOAX??)) two aliens appeared from a bright light to peacefully investigate the surroundings in the woods
2 Witnessed two aliens walking along baseball field fence.
```

'Find' functions: `str_subset()` is easier

`str_subset()` gives the values that match the pattern:

```
str_subset(ufo$comments, "two aliens")
```

```
[1] "((HOAX??))  two aliens appeared from a bright light to peacefully investigate the surroundings in the woods"  
[2] "Witnessed two aliens walking along baseball field fence."
```

Showing difference in `str_extract`

`str_extract` extracts just the matched string

```
ss = str_extract(ufo$comments, "two aliens")  
head(ss)
```

```
[1] NA NA NA NA NA NA
```

```
ss[ !is.na(ss) ]
```

```
[1] "two aliens" "two aliens"
```

- Look for any comment that starts with “aliens”

```
str_subset(ufo$comments, "^aliens.*")
```

```
[1] "aliens speak german???" "aliens exist"           "aliens in srilanka"
```

Using Regular Expressions

That contains space then ship maybe with stuff in between

```
str_subset(ufo$comments, "space.?ship") %>% head(4) # gets "spaceship" or "space ship" or...
```

```
[1] "I saw the cylinder shaped looked like a spaceship hovring above the east side of the Air Force base. Saw it for
[2] "description of a spaceship spotted over Birmingham Alabama in 1967."
[3] "A space ship was descending to the ground"
[4] "On Monday october 3&#44 2005&#44 I spotted two spaceships in the sky. The first spotted ship was what seemed t
```

```
str_subset(ufo$comments, "space.ship") %>% head(4) # no "spaceship" must have character in bw
```

```
[1] "A space ship was descending to the ground"
[2] "I saw a Silver space ship rising into the early morning sky over Houston##44 Texas."
[3] "Saw a space ship hanging over the southern (Manzano) portion of the Sandia Mountains on evening. It was brightl
[4] "saw space ship for 5 min##33 Got scared crapless##33##33##33##33##33##33##33##33##33##33##33##33##33##33##33##33
```

str_replace()

Let's say we wanted to make the time information more consistent. Using `case_when()` would be very tedious and error-prone!

We can use `str_replace()` to do so.

```
head(ufo$duration_hours_min, 8)
```

```
[1] "45 minutes"    "1-2 hrs"       "20 seconds"    "1/2 hour"      "15 minutes"
[8] "20 minutes"
```

```
ufo %>% mutate(duration_hours_min =
  str_replace(string = duration_hours_min,
    pattern = "minutes",
    replacement = "mins")) %>%
  pull(duration_hours_min) %>%
  head(8)
```

```
[1] "45 mins"      "1-2 hrs"       "20 seconds"    "1/2 hour"      "15 mins"
[8] "20 mins"
```

Dates and times

The [lubridate](<https://lubridate.tidyverse.org/>) package is amazing, there's no reason to use anything else.

```
library(lubridate) #need to load this one!
head(ufo$datetime)
```

```
[1] "10/10/1949 20:30" "10/10/1949 21:00" "10/10/1955 17:00" "10/10/1956 21:00" "10/10/1960 20:00" "10/10/1961 19:00"
```

```
ufo$date_posted = mdy(ufo$date_posted)
```

```
Warning: 194 failed to parse.
```

```
head(ufo$date_posted)
```

```
[1] "2004-04-27" "2005-12-16" "2008-01-21" "2004-01-17" "2004-01-22" "2007-04-27"
```