

Data Cleaning Part 2

Data Wrangling in R

Data Cleaning Part 2

Example of Cleaning: more complicated

For example, let's say gender was coded as Male, M, m, Female, F, f. Using Excel to find all of these would be a matter of filtering and changing all by hand or using if statements.

Sometimes though, it's not so simple. That's where functions that find patterns come to be very useful.

```
table(gender)
```

```
gender
```

F	FeMAle	FEMALE	Fm	M	Ma	mAle	Male	M
80	88	76	87	99	76	84	83	

```
Woman
```

```
71
```

Example of Cleaning: more complicated

In R, you could use `case_when()`:

```
#case_when way:  
data_gen <- data_gen %>% mutate(gender =  
  case_when(gender %in% c("Male", "M",  
    ~ "Male",  
    TRUE ~ gender))  
head(data_gen)
```

```
# A tibble: 6 x 1
```

```
  gender
```

```
  <chr>
```

```
1 F
```

```
2 Fm
```

```
3 MaLe
```

```
4 MaLe
```

```
5 FeMAle
```

```
6 FEMALE
```

Oh dear! This only fixes some values! It is difficult to notice values

String functions

The stringr package

Like dplyr, the stringr package:

- ▶ Makes some things more intuitive
- ▶ Is different than base R
- ▶ Is used on forums for answers
- ▶ Has a standard format for most functions: `str_`
 - ▶ the first argument is a string like first argument is a `data.frame` in dplyr

Useful String Functions

Useful String functions from base R and stringr

- ▶ `toupper()`, `tolower()` - uppercase or lowercase your data
- ▶ `str_sentence()` - uppercase just the first character (in the stringr package)
- ▶ `paste()` - paste strings together with a space
- ▶ `paste0` - paste strings together with no space as default
- ▶ `str_trim()` (in the stringr package) or `trimws` in base
 - ▶ will trim whitespace
- ▶ `nchar` - get the number of characters in a string

recoding with str_to_sentence()

```
#case_when way:  
data_gen <- data_gen %>%  
  mutate(gender = str_to_sentence(gender)) %>%  
  mutate(gender =  
    case_when(gender %in% c("Male", "M",  
      ~ "Male",  
      TRUE ~ gender))  
head(data_gen)
```

```
# A tibble: 6 x 1
```

```
  gender  
  <chr>
```

```
1 F  
2 Fm  
3 Male  
4 Male  
5 Female  
6 Female
```


Reading in again

Now we have a chance to keep but clean these values!

```
ufo <- read_csv("https://sisbid.github.io/Data-Wrangling/data/ufo.csv")
```

Warning: One or more parsing issues, call `problems()` on your object to see details.
e.g.:

```
dat <- vroom(...)
problems(dat)
```

```
p <- problems(ufo)
ufo_clean <- ufo %>% slice((pull(p, row))*-1)
```

Clean names with the `clean_names()` function from the `janitor` package

```
colnames(ufo_clean)
```

```
[1] "datetime"           "city"               "state"  
[4] "country"            "shape"              "duration"  
[7] "duration (hours/min)" "comments"           "date posted"  
[10] "latitude"           "longitude"
```

```
ufo_clean <- clean_names(ufo_clean)  
colnames(ufo_clean)
```

```
[1] "datetime"           "city"               "state"  
[4] "country"            "shape"              "duration_seconds"  
[7] "duration_hours_min" "comments"           "date_posted"  
[10] "latitude"           "longitude"
```

str_detect and filter

Now let's fix our ufo data and remove those pesky backticks in the duration_seconds variable. First let's find them with str_detect.

```
ufo_clean %>% filter(str_detect(string = duration_seconds,
```

```
# A tibble: 3 x 11
  datetime      city state country shape durat~1 durat~2 co
  <chr>         <chr> <chr> <chr>   <chr> <chr>   <chr>   <chr>
1 2/2/2000 19~ bouse az    us    <NA> 2`     each a~ Dr
2 4/10/2005 2~ sant~ ca    us    <NA> 8`     eight ~ 2
3 7/21/2006 1~ ibag~ <NA> <NA>   circ~ 0.5`  1/2 se~ V
# ... with 1 more variable: longitude <chr>, and abbreviated
# 1: duration_seconds, 2: duration_hours_min, 3: comments
# 5: latitude
```

str_remove

```
ufo_clean <- ufo_clean %>% mutate(duration_seconds = str_remove(duration_seconds, "[^0-9]"))

ufo_clean <- ufo_clean %>% mutate(duration_seconds = as.numeric(duration_seconds))
ufo_clean %>% filter(str_detect(string = duration_seconds, "[^0-9]"))

# A tibble: 0 x 11
# ... with 11 variables: datetime <chr>, city <chr>, state <chr>,
#   country <chr>, shape <chr>, duration_seconds <dbl>,
#   duration_hours_min <chr>, comments <chr>, date_posted <chr>,
#   latitude <chr>, longitude <chr>
```

Paste can add things back to variables

```
ufo_clean %>% mutate(duration_seconds = paste(duration_seconds,
```

```
Rows: 88,679
```

```
Columns: 11
```

```
$ datetime      <chr> "10/10/1949 20:30", "10/10/1949  
$ city          <chr> "san marcos", "lackland afb", "  
$ state         <chr> "tx", "tx", NA, "tx", "hi", "tn"  
$ country       <chr> "us", NA, "gb", "us", "us", "us"  
$ shape         <chr> "cylinder", "light", "circle", "  
$ duration_seconds <chr> "2700 sec", "7200 sec", "20 sec"  
$ duration_hours_min <chr> "45 minutes", "1-2 hrs", "20 sec"  
$ comments      <chr> "This event took place in early  
$ date_posted    <chr> "4/27/2004", "12/16/2005", "1/21  
$ latitude       <chr> "29.8830556", "29.38421", "53.2"  
$ longitude      <chr> "-97.9411111", "-98.581082", "-2"
```

Substringing

stringr

- ▶ `str_sub(x, start, end)` - substrings from position start to position end
- ▶ `str_split(string, pattern)` - splits strings up - returns list! [we'll revisit in "Functional Programming"]

Substringing

Examples:

```
str_sub("I like friesian horses", 8,12)
```

```
[1] "fries"
```

```
#123456789101112
```

```
#I like fries
```

```
str_sub(c("Site A", "Site B", "Site C"), 6,6)
```

```
[1] "A" "B" "C"
```

Splitting/Find/Replace and Regular Expressions

- ▶ R can do much more than find exact matches for a whole string
- ▶ Like Perl and other languages, it can use regular expressions.
- ▶ What are regular expressions?
 - ▶ Ways to search for specific strings
 - ▶ Can be very complicated or simple
 - ▶ Highly Useful - think “Find” on steroids

A bit on Regular Expressions

- ▶ <http://www.regular-expressions.info/reference.html>
- ▶ They can use to match a large number of strings in one statement
- ▶ . matches any single character
- ▶ * means repeat as many (even if 0) more times the last character
- ▶ ? makes the last thing optional
- ▶ ^ matches start of vector ^a - starts with "a"
- ▶ \$ matches end of vector b\$ - ends with "b"

'Find' functions: `stringr`

`str_detect`, `str_subset`, `str_replace`, and `str_replace_all` search for matches to argument pattern within each element of a character vector: they differ in the format of and amount of detail in the results.

- ▶ `str_detect` - returns TRUE if pattern is found
- ▶ `str_subset` - returns only the strings which pattern were detected
 - ▶ convenient wrapper around `x[str_detect(x, pattern)]`
- ▶ `str_extract` - returns only strings which pattern were detected, but ONLY the pattern
- ▶ `str_replace` - replaces pattern with replacement the first time
- ▶ `str_replace_all` - replaces pattern with replacement as many times matched

'Find' functions: Finding Indices

These are the indices where the pattern match occurs:

```
ufo_clean %>% filter(str_detect(comments, "two aliens")) %>%
```

```
# A tibble: 2 x 11
```

	datetime	city	state	country	shape	durat~1	durat~2	co
	<chr>	<chr>	<chr>	<chr>	<chr>	<dbl>	<chr>	<chr>
1	10/14/2006	~ yuma	va	us	form~	300	5 minu~	(
2	7/1/2007	23~ nort~	ct	<NA>	unkn~	60	1 minu~	W

```
# ... with 1 more variable: longitude <chr>, and abbreviated  
# 1: duration_seconds, 2: duration_hours_min, 3: comments  
# 5: latitude
```

```
ufo_clean %>% filter(str_detect(comments, "two aliens")) %>%
```

```
# A tibble: 1 x 1
```

	n
	<int>
1	2

'Find' functions: Finding Logicals

filter() using str_detect() gives a tibble:

```
filter(ufo_clean, str_detect(comments, "two aliens"))
```

```
# A tibble: 2 x 11
```

	datetime	city	state	country	shape	durat~1	durat~2	co
	<chr>	<chr>	<chr>	<chr>	<chr>	<dbl>	<chr>	<chr>
1	10/14/2006	~ yuma	va	us	form~	300	5 minu~	(
2	7/1/2007	23~ nort~	ct	<NA>	unkn~	60	1 minu~	W

```
# ... with 1 more variable: longitude <chr>, and abbreviated
# 1: duration_seconds, 2: duration_hours_min, 3: comments
# 5: latitude
```

```
filter(ufo_clean, str_detect(comments, "two aliens")) %>% s
```

```
# A tibble: 2 x 1
```

	comments
	<chr>
1	((HOAX??)) two aliens appeared from a bright light to pe
2	Witnessed two aliens walking along baseball field fence.

'Find' functions: `str_subset()` is easier

`str_subset()` gives the values that match the pattern:

```
ufo_clean %>% pull(comments) %>% str_subset("two aliens")
```

```
[1] "((HOAX??))  two aliens appeared from a bright light to  
[2] "Witnessed two aliens walking along baseball field fence"
```

Showing difference in str_extract

str_extract extracts just the matched string

```
ufo_clean %>% mutate(two = str_extract(comments, "two")) %>%
```

```
# A tibble: 2 x 2
```

```
  two      n
```

```
  <chr> <int>
```

```
1 two      2497
```

```
2 <NA>    86182
```

- ▶ Look for any comment that starts with "aliens"

```
ufo_clean %>% pull(comments) %>% str_subset( "^aliens.*")
```

```
[1] "aliens speak german???" "aliens exist" "aliens exist"
```

Using Regular Expressions

That contains space then ship maybe with stuff in between

```
ufo_clean %>% pull(comments) %>% str_subset("space.?ship") %>%
```

```
[1] "I saw the cylinder shaped looked like a spaceship hovering over  
[2] "description of a spaceship spotted over Birmingham Alabama  
[3] "A space ship was descending to the ground"  
[4] "On Monday october 3rd 2005 I spotted two spaceships flying over"
```

```
ufo_clean %>% pull(comments) %>% str_subset("space.ship") %>%
```

```
[1] "A space ship was descending to the ground"  
[2] "I saw a Silver space ship rising into the early morning sky  
[3] "Saw a space ship hanging over the southern (Manzano) part of  
[4] "saw space ship for 5 min Got scared crapless"
```

str_replace()

Let's say we wanted to make the time information more consistent. Using `case_when()` would be very tedious and error-prone!

We can use `str_replace()` to do so.

```
head(pull(ufo_clean, duration_hours_min))
```

```
[1] "45 minutes" "1-2 hrs"      "20 seconds" "1/2 hour"    "15  
[6] "5 minutes"
```

```
ufo_clean %>% mutate(duration_hours_min =  
                      str_replace(string = duration_hours_min,  
                                   pattern = "minutes",  
                                   replacement = "mins")) %>%  
  pull(duration_hours_min) %>%  
  head(8)
```

```
[1] "45 mins"      "1-2 hrs"      "20 seconds"   "1/2 hour"  
[6] "5 mins"       "about 3 mins" "20 mins"
```


Separating columns

Better yet, you might notice that this data isn't tidy- there are more than two entries for each value - amount of time and unit. We could separate this using `separate()` from the `tidyr` package.

```
ufo_clean %>% separate(duration_hours_min,  
                        into = c("duration_amount", "duration_unit"),  
                        sep = " ") %>%  
  select(duration_amount, duration_unit) %>% head()
```

```
# A tibble: 6 x 2  
  duration_amount duration_unit  
    <chr>          <chr>  
1 45             minutes  
2 1-2            hrs  
3 20             seconds  
4 1/2            hour  
5 15             minutes  
6 5              minutes
```

As you can see there is still plenty of cleaning to do!

Dates and times

The [lubridate](https://lubridate.tidyverse.org/) package is amazing for dates. Most important functions are those that look like ymd or mdy etc. They specify how a date should be interpreted.

```
library(lubridate)#need to load this one!
```

```
ufo_clean <- ufo_clean %>% mutate(date_posted = mdy(date_po
```

Warning: 193 failed to parse.

```
head(ufo_clean)
```

```
# A tibble: 6 x 11
```

	datetime	city	state	country	shape	durat~1	durat~2	comme
	<chr>	<chr>	<chr>	<chr>	<chr>	<dbl>	<chr>	<chr>
1	10/10/19~	san ~	tx	us	cyli~	2700	45 min~	This
2	10/10/19~	lack~	tx	<NA>	light	7200	1-2 hrs	1949
3	10/10/19~	ches~	<NA>	gb	circ~	20	20 sec~	Green
4	10/10/19~	edna	tx	us	circ~	20	1/2 ho~	My o
5	10/10/19~	kane~	hi	us	light	900	15 min~	AS a

Summary

- ▶ `stringr` package has lots of helpful functions that work on vectors or variables in a data frame
- ▶ `str_detect` helps find patterns
- ▶ `str_detect` and `filter` can help you filter data based on patterns within value
- ▶ `str_extract` helps extract a pattern
- ▶ `str_sub` extracts pieces of strings based on the position of the the characters
- ▶ `str_subset` gives the values that match a pattern
- ▶ `separate` can separate columns into two
- ▶ `^` indicates the start of a string
- ▶ `$` indicates the end of a string
- ▶ the `lubridate` package is useful for dates and times

Lab

<https://sisbid.github.io/Data-Wrangling/labs/data-cleaning-lab-part2.Rmd>