

Data I/O + Structure

Data Wrangling in R

What did I just read in?

- `nrow()` displays the number of rows of a data frame
- `ncol()` displays the number of columns
- `dim()` displays a vector of length 2: # rows, # columns

```
nrow(ufo)
```

```
[1] 88875
```

```
ncol(ufo)
```

```
[1] 11
```

```
dim(ufo)
```

```
[1] 88875    11
```

All Column Names

- `colnames()` displays the column names

```
colnames(ufo)
```

```
[1] "datetime"      "city"           "state"
[4] "country"       "shape"          "duration (seconds)"
[7] "duration (hours/min)" "comments"       "date posted"
[10] "latitude"      "longitude"
```

Structure using `str()`

```
str(ufo)
```

```
spec_tbl_df [88,875 × 11] (S3: spec_tbl_df/tbl_df/tbl/data.frame)
 $ datetime      : chr [1:88875] "10/10/1949 20:30" "10/10/1949 21:00" "
 $ city          : chr [1:88875] "san marcos" "lackland afb" "chester (u
 $ state         : chr [1:88875] "tx" "tx" NA "tx" ...
 $ country       : chr [1:88875] "us" NA "gb" "us" ...
 $ shape         : chr [1:88875] "cylinder" "light" "circle" "circle" ..
 $ duration (seconds) : num [1:88875] 2700 7200 20 20 900 300 180 1200 180 12
 $ duration (hours/min): chr [1:88875] "45 minutes" "1-2 hrs" "20 seconds" "1/
 $ comments      : chr [1:88875] "This event took place in early fall ar
 $ date posted   : chr [1:88875] "4/27/2004" "12/16/2005" "1/21/2008" "1
 $ latitude      : chr [1:88875] "29.8830556" "29.38421" "53.2" "28.9783
 $ longitude     : chr [1:88875] "-97.9411111" "-98.581082" "-2.916667"
 - attr(*, "spec")=
  .. cols(
  ..   datetime = col_character(),
  ..   city = col_character(),
  ..   state = col_character(),
  ..   country = col_character(),
  ..   shape = col_character(),
  ..   `duration (seconds)` = col_double(),
  ..   `duration (hours/min)` = col_character(),
  ..   comments = col_character(),
  ..   `date posted` = col_character(),
  ..   latitude = col_character(),
  ..   longitude = col_character()
```

Data Input

- Sometimes you get weird messages when reading in data.
- The `problems()` function shows you any issues with the data read-in.

```
head(problems(ufo))
```

```
# A tibble: 6 × 5
  row   col expected actual      file
<int> <int> <chr>      <chr>    <chr>
1   878    12 11 columns 12 columns /Users/avahoffman/Dropbox/JHSPH/Data-Wrang
2  1713    12 11 columns 12 columns /Users/avahoffman/Dropbox/JHSPH/Data-Wrang
3  1815    12 11 columns 12 columns /Users/avahoffman/Dropbox/JHSPH/Data-Wrang
4  2858    12 11 columns 12 columns /Users/avahoffman/Dropbox/JHSPH/Data-Wrang
5  3734    12 11 columns 12 columns /Users/avahoffman/Dropbox/JHSPH/Data-Wrang
6  4756    12 11 columns 12 columns /Users/avahoffman/Dropbox/JHSPH/Data-Wrang
```


```
dim(problems(ufo))
```

```
[1] 199    5
```

Missing data with the **vizdat**/**naniar** packages


The `vis_dat` function can give you an overview

```
library(vizdat)
ufo_samp <- ufo %>% sample_n(size = 1000) # Subset for big data
vis_dat(ufo_samp)
```




Missing data with the **vizdat**/**naniar** packages

```
vis_miss(ufo_samp)
```



Missing data with the `vizdat`/`naniar` packages

```
library(naniar)  
gg_miss_var(ufo)
```



Missing data with the **vizdat**/**naniar** packages

`miss_case_summary` which rows have missing data in order

```
miss_case_summary(ufo)
```

```
# A tibble: 88,875 × 3
  case n_miss pct_miss
  <int>   <int>   <dbl>
1     877     6     54.5
2    1712     6     54.5
3    5613     6     54.5
4    7515     6     54.5
5    7625     6     54.5
6   10156     6     54.5
7   16634     6     54.5
8   18240     6     54.5
9   19813     6     54.5
10  19908     6     54.5
# ... with 88,865 more rows
```

Missing data with the **vizdat**/**naniar** packages

`miss_var_summary` which variables have missing data

```
miss_var_summary(ufo)
```

```
# A tibble: 11 × 3
  variable          n_miss pct_miss
  <chr>          <int>    <dbl>
1 country       12561  14.1
2 state         7519   8.46
3 shape         3118   3.51
4 duration (hours/min) 3101   3.49
5 city          196   0.221
6 comments      126   0.142
7 duration (seconds)    5  0.00563
8 datetime        0    0
9 date posted         0    0
10 latitude        0    0
11 longitude        0    0
```

After hours of cleaning...

More ways to save: write_rds

If you want to save **one** object, you can use `readr::write_rds` to save to a compressed `rds` file:

```
write_rds(ufo, file = "ufo_dataset.rds", compress = "xz")
```

More ways to save: `read_rds`

To read this back in to R, you need to use `read_rds`, but **need to assign it**:

```
ufo3 <- read_rds(file = "ufo_dataset.rds")  
identical(ufo, ufo3) # test if they are the same
```

```
[1] FALSE
```

More ways to save: **save**

The `save` command can save a set of R objects into an “R data file”, with the extension `.rda` or `.RData`.

```
x = 5  
save(ufo, x, file = "ufo_data.rda")
```

More ways to save: load

The opposite of `save` is `load`. The `ls()` command lists the items in the workspace/environment and `rm()` removes them

```
load(file = "ufo_data.rda")
```

Data Output

While its nice to be able to read in a variety of data formats, it's equally important to be able to output data somewhere.

`write_delim()`: Write a data frame to a delimited file "This is about twice as fast as `write.csv()`, and never writes row names."

Data Output

`x`: A data frame to write to disk

`file`: the file name where you want to R object written. It can be an absolute path, or a filename (which writes the file to your working directory)

`delim`: what character separates the columns?

- `","` = .csv - Note there is also a `write_csv()` function
- `"\t"` = tab delimited

Data Output

For example, we can write back out just the first 100 lines of the `ufo` dataset:

```
write_delim(ufo[1:100,], file = "ufo_first100.csv", delim = ",")
```