

Data Reshaping

Data Wrangling in R

Reshaping Data

In this module, we will show you how to:

1. Reshape data from wide (fat) to long (tall)
2. Reshape data from long (tall) to wide (fat)
3. Merge Data/Joins
4. Perform operations by a grouping variable

What is wide/long data?

Data is stored *differently* in the tibble.

Wide: has many columns

```
# A tibble: 1 x 4
  State    June_vacc_rate May_vacc_rate April_vacc_rate
<chr>    <chr>          <chr>          <chr>
1 Alabama 37.2%          36.0%          32.4%
```

Long: column names become data

```
# A tibble: 3 x 3
  State    name          value
<chr>    <chr>          <chr>
1 Alabama June_vacc_rate 37.2%
2 Alabama May_vacc_rate  36.0%
3 Alabama April_vacc_rate 32.4%
```

What is wide/long data?

Wide: multiple columns per individual, values spread across multiple columns

```
# A tibble: 2 x 4
  State      June_vacc_rate May_vacc_rate April_vacc_rate
  <chr>      <chr>           <chr>         <chr>
1 Alabama  37.2%              36.0%         32.4%
2 Alaska   47.5%              46.2%         41.7%
```

Long: multiple rows per observation, a single column contains the values

```
# A tibble: 6 x 3
  State      name          value
  <chr>      <chr>           <chr>
1 Alabama  June_vacc_rate  37.2%
2 Alabama  May_vacc_rate   36.0%
3 Alabama  April_vacc_rate 32.4%
4 Alaska   June_vacc_rate  47.5%
5 Alaska   May_vacc_rate   46.2%
6 Alaska   April_vacc_rate 41.7%
```

What is wide/long data?

Data is wide or long **with respect** to certain variables.

	Day 1	Day 2	Day 3
Patient 1	A	B	C
Patient 2	D	E	F

Wide

Long

	Day	Value
Patient 1	Day 1	A
Patient 1	Day 2	B
Patient 1	Day 3	C
Patient 2	Day 1	D
Patient 2	Day 2	E
Patient 2	Day 3	F

CC-BY jhudatascience.org

What is wide/long data?

wide

	wide		
id	x	y	z
1	a	c	e
2	b	d	f

<https://github.com/gadenbuie/tidyexplain/blob/master/images/tidyr-spread-gather.gif>

Why do we need to switch between wide/long data?

Wide: Easier for humans to read

```
# A tibble: 2 x 4
  State      June_vacc_rate May_vacc_rate April_vacc_rate
  <chr>      <chr>           <chr>         <chr>
1 Alabama  37.2%             36.0%         32.4%
2 Alaska   47.5%             46.2%         41.7%
```

Long: Easier for R to make plots & do analysis

```
# A tibble: 6 x 3
  State      name      value
  <chr>      <chr>      <chr>
1 Alabama June_vacc_rate 37.2%
2 Alabama May_vacc_rate  36.0%
3 Alabama April_vacc_rate 32.4%
4 Alaska  June_vacc_rate 47.5%
5 Alaska  May_vacc_rate  46.2%
6 Alaska  April_vacc_rate 41.7%
```

Data used: Charm City Circulator

http://jhudatascience.org/intro_to_r/data/Charm_City_Circulator_Ridership.csv

```
circ = read_csv(  
  paste0("http://jhudatascience.org/intro_to_r/",  
        "data/Charm_City_Circulator_Ridership.csv"))  
head(circ, 5)
```

```
# A tibble: 5 x 15  
  day      date      orangeBoardings orangeAlightings orangeAverage purpleBoardin  
  <chr>    <chr>          <dbl>             <dbl>             <dbl>          <dbl>  
1 Monday  01/11/...         877             1027             952  
2 Tuesday 01/12/...         777             815             796  
3 Wednes... 01/13/...        1203            1220            1212.  
4 Thursd... 01/14/...        1194            1233            1214.  
5 Friday   01/15/...        1645            1643            1644  
# ... with 9 more variables: purpleAlightings <dbl>, purpleAverage <dbl>,  
#   greenBoardings <dbl>, greenAlightings <dbl>, greenAverage <dbl>,  
#   bannerBoardings <dbl>, bannerAlightings <dbl>, bannerAverage <dbl>,  
#   daily <dbl>
```


tidyr package

`tidyr` allows you to “tidy” your data. We will be talking about:

- `pivot_longer` - make multiple columns into variables, (wide to long)
- `pivot_wider` - make a variable into multiple columns, (long to wide)
- `separate` - string into multiple columns
- `unite` - multiple columns into one string

The `reshape` command exists. It is a **confusing** function. Don't use it.

Reshaping data from wide (fat) to long (tall): tidyr

`tidyr::pivot_longer` - puts column data into rows.

- First describe which columns we want to “pivot_longer”
- `names_to` = gives a new name to the pivoted columns
- `values_to` = gives a new name to the values that used to be in those columns

```
long = circ %>%  
  pivot_longer(starts_with(c("orange", "purple", "green", "banner")),  
               names_to = "var", values_to = "number")  
long
```

```
# A tibble: 13,752 x 5  
   day   date      daily var          number  
   <chr> <chr>    <dbl> <chr>    <dbl>  
1 Monday 01/11/2010    952 orangeBoardings    877  
2 Monday 01/11/2010    952 orangeAlightings  1027  
3 Monday 01/11/2010    952 orangeAverage     952  
4 Monday 01/11/2010    952 purpleBoardings    NA  
5 Monday 01/11/2010    952 purpleAlightings    NA  
6 Monday 01/11/2010    952 purpleAverage      NA  
7 Monday 01/11/2010    952 greenBoardings     NA  
8 Monday 01/11/2010    952 greenAlightings     NA  
9 Monday 01/11/2010    952 greenAverage       NA
```

Reshaping data from wide (fat) to long (tall): tidyr

We have many columns here, so we could instead use the `!` to say which columns we *don't* want to pivot.

```
long = circ %>% pivot_longer(!c(day, date, daily),  
                             names_to = "var", values_to = "number")  
long
```

```
# A tibble: 13,752 x 5  
   day      date      daily var              number  
   <chr>   <chr>    <dbl> <chr>              <dbl>  
1 Monday 01/11/2010    952 orangeBoardings     877  
2 Monday 01/11/2010    952 orangeAlightings   1027  
3 Monday 01/11/2010    952 orangeAverage       952  
4 Monday 01/11/2010    952 purpleBoardings     NA  
5 Monday 01/11/2010    952 purpleAlightings    NA  
6 Monday 01/11/2010    952 purpleAverage       NA  
7 Monday 01/11/2010    952 greenBoardings      NA  
8 Monday 01/11/2010    952 greenAlightings     NA  
9 Monday 01/11/2010    952 greenAverage        NA  
10 Monday 01/11/2010    952 bannerBoardings     NA  
# ... with 13,742 more rows
```

Reshaping data from wide (fat) to long (tall): tidyr

```
long %>% count(var)
```

```
# A tibble: 12 x 2
```

	var <chr>	n <int>
1	bannerAlightings	1146
2	bannerAverage	1146
3	bannerBoardings	1146
4	greenAlightings	1146
5	greenAverage	1146
6	greenBoardings	1146
7	orangeAlightings	1146
8	orangeAverage	1146
9	orangeBoardings	1146
10	purpleAlightings	1146
11	purpleAverage	1146
12	purpleBoardings	1146

Making a separator

We will use `str_replace` from the `stringr` package to put `_` in the names

```
long = long %>% mutate(  
  var = str_replace(var, "Board", " _Board"),  
  var = str_replace(var, "Alight", " _Alight"),  
  var = str_replace(var, "Average", " _Average")  
)  
long
```

```
# A tibble: 13,752 x 5  
  day    date      daily var              number  
  <chr> <chr>    <dbl> <chr>              <dbl>  
1 Monday 01/11/2010    952 orange_Boardings      877  
2 Monday 01/11/2010    952 orange_Alightings    1027  
3 Monday 01/11/2010    952 orange_Average       952  
4 Monday 01/11/2010    952 purple_Boardings      NA  
5 Monday 01/11/2010    952 purple_Alightings     NA  
6 Monday 01/11/2010    952 purple_Average       NA  
7 Monday 01/11/2010    952 green_Boardings      NA  
8 Monday 01/11/2010    952 green_Alightings     NA  
9 Monday 01/11/2010    952 green_Average       NA  
10 Monday 01/11/2010    952 banner_Boardings     NA  
# ... with 13,742 more rows
```

Reshaping data from wide (fat) to long (tall): tidyr

Now each `var` is Boardings, Averages, or Alightings. We use `"into ="` to name the new columns and `"sep ="` to show where the separation should happen.

```
long =  
  long %>%  
    separate(var, into = c("line", "type"), sep = "_")  
long
```

```
# A tibble: 13,752 x 6  
  day    date      daily line    type      number  
  <chr> <chr>    <dbl> <chr> <chr>    <dbl>  
1 Monday 01/11/2010    952 orange Boardings    877  
2 Monday 01/11/2010    952 orange Alightings  1027  
3 Monday 01/11/2010    952 orange Average     952  
4 Monday 01/11/2010    952 purple Boardings     NA  
5 Monday 01/11/2010    952 purple Alightings     NA  
6 Monday 01/11/2010    952 purple Average     NA  
7 Monday 01/11/2010    952 green Boardings     NA  
8 Monday 01/11/2010    952 green Alightings     NA  
9 Monday 01/11/2010    952 green Average     NA  
10 Monday 01/11/2010    952 banner Boardings     NA  
# ... with 13,742 more rows
```

Re-uniting all the lines

If we had the opposite problem, we could use the `unite` function:

```
reunited = long %>%  
  unite(var, line, type, sep = "_")  
reunited
```

```
# A tibble: 13,752 x 5  
  day      date      daily var      number  
  <chr>   <chr>    <dbl> <chr>    <dbl>  
1 Monday 01/11/2010    952 orange_Boardings    877  
2 Monday 01/11/2010    952 orange_Alightings  1027  
3 Monday 01/11/2010    952 orange_Average     952  
4 Monday 01/11/2010    952 purple_Boardings    NA  
5 Monday 01/11/2010    952 purple_Alightings    NA  
6 Monday 01/11/2010    952 purple_Average      NA  
7 Monday 01/11/2010    952 green_Boardings     NA  
8 Monday 01/11/2010    952 green_Alightings    NA  
9 Monday 01/11/2010    952 green_Average       NA  
10 Monday 01/11/2010    952 banner_Boardings    NA  
# ... with 13,742 more rows
```

Reshaping data from long (tall) to wide (fat): tidyr

In `tidyr`, the `pivot_wider` function spreads rows into columns. Now we have a long data set, but we want to separate the Average, Alightings and Boardings into different columns:

```
wide = long %>% pivot_wider(names_from = "type",  
                             values_from = "number")
```

```
wide
```

```
# A tibble: 4,584 x 7  
  day      date      daily line      Boardings Alightings Average  
  <chr>    <chr>    <dbl> <chr>    <dbl>      <dbl>    <dbl>  
1 Monday  01/11/2010  952 orange      877        1027      952  
2 Monday  01/11/2010  952 purple       NA         NA       NA  
3 Monday  01/11/2010  952 green        NA         NA       NA  
4 Monday  01/11/2010  952 banner       NA         NA       NA  
5 Tuesday 01/12/2010  796 orange      777         815      796  
6 Tuesday 01/12/2010  796 purple       NA         NA       NA  
7 Tuesday 01/12/2010  796 green        NA         NA       NA  
8 Tuesday 01/12/2010  796 banner       NA         NA       NA  
9 Wednesday 01/13/2010 1212. orange    1203        1220    1212.  
10 Wednesday 01/13/2010 1212. purple       NA         NA       NA  
# ... with 4,574 more rows
```


`pivot_longer`

```
sandwich %>%  
  pivot_longer()
```

