

Subsetting Data in R

Data Wrangling in R

UFO data again

```
ufo <- read_csv("../data/ufo/ufo_data_complete.csv", col_types =  
  cols(  
    .default = col_character(),  
    `duration (seconds)` = col_double(),  
    longitude = col_double()  
  ))  
  
head(ufo)
```

```
# A tibble: 6 x 11
```

| | datetime | city | state | country | shape | durat~1 | durat~2 | co |
|---|------------|---------|-------|---------|-------|---------|---------|-------|
| | <chr> | <chr> | <chr> | <chr> | <chr> | <dbl> | <chr> | <chr> |
| 1 | 10/10/1949 | ~ san | ~ tx | us | cyli~ | 2700 | 45 min~ | TH |
| 2 | 10/10/1949 | ~ lack~ | tx | <NA> | light | 7200 | 1-2 hrs | 19 |
| 3 | 10/10/1955 | ~ ches~ | <NA> | gb | circ~ | 20 | 20 sec~ | Gr |
| 4 | 10/10/1956 | ~ edna | tx | us | circ~ | 20 | 1/2 ho~ | My |
| 5 | 10/10/1960 | ~ kane~ | hi | us | light | 900 | 15 min~ | AS |
| 6 | 10/10/1961 | ~ bris~ | tn | us | sphe~ | 300 | 5 minu~ | My |

```
# ... with 1 more variable: longitude <dbl>, and abbreviated  
# 1: `duration (seconds)` 2: `duration (hours/min)` 3:
```

Filtering and tibbles

The filter function automatically removes NA values.

```
count(ufo, country)
```

```
# A tibble: 6 x 2
  country      n
  <chr>    <int>
1 au         593
2 ca        3266
3 de         112
4 gb        2050
5 us       70293
6 <NA>     12561
```

```
ufo %>%filter(country == "de") %>% count(country)
```

```
# A tibble: 1 x 2
  country      n
  <chr>    <int>
1 de         112
```

Filtering for missing data

Missing value and filter can be powerful

```
ufo %>%  
  filter(is.na(state) & is.na(country)) %>%  
  head()
```

```
# A tibble: 6 x 11
```

| | datetime | city | state | country | shape | durat~1 | durat~2 | co |
|---|------------|---------|-------|---------|-------|---------|---------|-------|
| | <chr> | <chr> | <chr> | <chr> | <chr> | <dbl> | <chr> | <chr> |
| 1 | 10/10/1973 | ~ berm~ | <NA> | <NA> | light | 20 | 20 sec. | sa |
| 2 | 10/10/1982 | ~ gisb~ | <NA> | <NA> | disk | 120 | 2min | gr |
| 3 | 10/10/1993 | ~ zlat~ | <NA> | <NA> | sphe~ | 1200 | 20 min~ | I |
| 4 | 10/10/1996 | ~ lake~ | <NA> | <NA> | light | 300 | 5 min | RE |
| 5 | 10/10/1998 | ~ turi~ | <NA> | <NA> | disk | 15 | 15 sec | We |
| 6 | 10/10/1999 | ~ dors~ | <NA> | <NA> | flash | 1 | <1 sec~ | WH |

```
# ... with 1 more variable: longitude <dbl>, and abbreviated  
# 1: `duration (seconds)`, 2: `duration (hours/min)`, 3:  
# 4: `date posted`, 5: latitude
```

Filtering and tibbles

Group logical statements with parentheses

```
ufo %>%  
  filter(  
    (!is.na(state) & is.na(country)) | city == "seattle") %>%  
  head()
```

```
# A tibble: 6 x 11
```

| | datetime | city | state | country | shape | durat~1 | durat~2 | co |
|---|------------|---------|-------|---------|-------|---------|---------|-------|
| | <chr> | <chr> | <chr> | <chr> | <chr> | <dbl> | <chr> | <chr> |
| 1 | 10/10/1949 | ~ lack~ | tx | <NA> | light | 7200 | 1-2 hrs | 19 |
| 2 | 10/10/1974 | ~ will~ | az | <NA> | light | 120 | 2 min | Th |
| 3 | 10/10/1979 | ~ sadd~ | ab | <NA> | tria~ | 270 | 4.5 or~ | Li |
| 4 | 10/10/1986 | ~ holm~ | ny | <NA> | chev~ | 180 | 3 minu~ | For |
| 5 | 10/10/1987 | ~ mani~ | on | <NA> | disk | 600 | 10/mins | We |
| 6 | 10/10/1989 | ~ kran~ | ky | <NA> | tria~ | 180 | 3min | Tr |

```
# ... with 1 more variable: longitude <dbl>, and abbreviated
```

```
# 1: `duration (seconds)`, 2: `duration (hours/min)`, 3:
```

```
# 4: `date posted`, 5: latitude
```

Renaming Columns

Renaming Columns of a data.frame

To rename columns in `dplyr`, you use the `rename` command (NEW=old)

```
ufo <- ufo %>% rename(City = city, duration_s = `duration`)
glimpse(ufo)
```

Rows: 88,875

Columns: 11

```
$ datetime      <chr> "10/10/1949 20:30", "10/10/1949 20:30"
```

```
$ City      <chr> "san marcos", "lackland afb"
```

```
$ state      <chr> "tx", "tx", NA, "tx", "hi",
```

```
$ country      <chr> "us", NA, "gb", "us", "us",
```

```
$ shape      <chr> "cylinder", "light", "circle"
```

```
$ duration_s      <dbl> 2700, 7200, 20, 20, 900, 300
```

```
$ `duration (hours/min)` <chr> "45 minutes", "1-2 hrs", "20"
```

```
$ comments      <chr> "This event took place in ea
```

```
$ `date posted`      <chr> "4/27/2004", "12/16/2005", "
```

```
$ latitude      <chr> "29.8830556", "29.38421", "50.00000"
```

```
$ longitude      <dbl> -97.941111, -98.581082, -2.9
```

Renaming All Columns of a data.frame: dplyr

To rename all columns you use the `rename_with` command (with a function)

```
ufo_upper <- ufo %>% rename_with(toupper)
glimpse(ufo_upper)
```

Rows: 88,875

Columns: 11

```
$ DATETIME      <chr> "10/10/1949 20:30", "10/10/1949 20:30"
```

```
$ CITY      <chr> "san marcos", "lackland afb"
```

```
$ STATE      <chr> "tx", "tx", NA, "tx", "hi",
```

```
$ COUNTRY      <chr> "us", NA, "gb", "us", "us",
```

```
$ SHAPE <chr> "cylinder", "light", "circle"
```

```
$ DURATION_S <dbl> 2700, 7200, 20, 20, 900, 300
```

```
$ `DURATION (HOURS/MIN)` <chr> "45 minutes", "1-2 hrs", "20"
```

```
$ COMMENTS      <chr> "This event took place in ea
```

```
$ `DATE POSTED`      <chr> "4/27/2004", "12/16/2005", "
```

```
$ LATITUDE      <chr> "29.8830556", "29.38421", "5
```

```
$ LONGITUDE      <dbl> -97.941111, -98.581082, -2.9
```


Adding columns to a data.frame

mutate - allows you to add or replace columns (IMPORTANT: need to reassign for it to stick!)

```
ufo2 <- ufo %>% mutate(State = toupper(state))
ufo2 %>% glimpse()
```

Rows: 88,875

Columns: 12

[illegible]

Creating conditional variables

One frequently-used tool is creating variables with conditions.

A general function for creating new variables based on existing variables is the `ifelse()` function, which “returns a value with the same shape as test which is filled with elements selected from either yes or no depending on whether the element of test is TRUE or FALSE.”

```
ifelse(test, yes, no)
```

```
# test: an object which can be coerced  
      to logical mode.
```

```
# yes: return values for true elements of test.
```

```
# no: return values for false elements of test.
```

Recoding to missing

Sometimes people code missing data in weird or inconsistent ways.

```
ages <- data.frame(age = c(23,-999,21,44,32,57,65,54))  
range(ages$age)
```

```
[1] -999    65
```

Recoding to create new column

Say we want to make a new column about if the age was over 30?

```
pull(ages, age)
```

```
[1] 23 -999 21 44 32 57 65 54
```

```
ages <- ages %>% mutate(over_20 = ifelse(age > 30, "Yes", "No"))
ages
```

| | age | over_20 |
|---|------|---------|
| 1 | 23 | No |
| 2 | -999 | No |
| 3 | 21 | No |
| 4 | 44 | Yes |
| 5 | 32 | Yes |
| 6 | 57 | Yes |
| 7 | 65 | Yes |
| 8 | 54 | Yes |

Recoding value to missing

How do we change the -999 to be treated as missing for the age column?

```
pull(ages, age)
```

```
[1] 23 -999 21 44 32 57 65 54
```

```
ages <- ages %>% mutate(age = ifelse(age == -999, NA, age))  
range(ages$age)
```

```
[1] NA NA
```

```
range(ages$age, na.rm=TRUE)
```

```
[1] 21 65
```

```
pull(ages, age)
```

```
[1] 23 NA 21 44 32 57 65 54
```

Adding columns to a data.frame: dplyr

```
ufo <- ufo %>% mutate(  
  region = ifelse(  
    country %in% c("us", "ca"),  
    "North America",  
    "Not North America")  
)  
ufo %>% select(country, region) %>% head()
```

A tibble: 6 x 2

| | country | region |
|---|---------|-------------------|
| | <chr> | <chr> |
| 1 | us | North America |
| 2 | <NA> | Not North America |
| 3 | gb | Not North America |
| 4 | us | North America |
| 5 | us | North America |
| 6 | us | North America |

case_when provides a more general way

```
casewhen(test ~ value if test is true,  
          test2 ~ vluue if test2 is true,  
          TRUE ~ value if all above tests are not true) # de
```

```
ufo <- ufo %>% mutate(  
  region = case_when(  
    country %in% c("us", "ca") ~ "North America",  
    country %in% c("de") ~ "Europe",  
    country %in% "gb" ~ "Great Britain",  
  ))
```

```
ufo %>% select(country, region) %>% head()
```

```
# A tibble: 6 x 2
```

| | country | region |
|---|---------|---------------|
| | <chr> | <chr> |
| 1 | us | North America |
| 2 | <NA> | <NA> |
| 3 | gb | Great Britain |

case_when defaults to NA when all tests fail

```
ufo <- ufo %>% mutate(  
  region = case_when(  
    country %in% c("us", "ca") ~ "North America",  
    country %in% c("de") ~ "Europe",  
    #country %in% "gb" ~ "Great Britain",  
  ))  
  
ufo %>% select(country, region) %>% head()
```

```
# A tibble: 6 x 2  
  country region  
  <chr>    <chr>  
1 us      North America  
2 <NA>    <NA>  
3 gb      <NA>  
4 us      North America  
5 us      North America  
6 us      North America
```


case_when() with value if all tests fail

```
ufo <- ufo %>% mutate(  
  region = case_when(  
    country %in% c("us", "ca") ~ "North America",  
    country %in% c("de") ~ "Europe",  
    country %in% "gb" ~ "Great Britain",  
    TRUE ~ "Other"  
  ))  
ufo %>% select(country, region) %>% head()
```

A tibble: 6 x 2

| | country | region |
|---|---------|---------------|
| | <chr> | <chr> |
| 1 | us | North America |
| 2 | <NA> | Other |
| 3 | gb | Great Britain |
| 4 | us | North America |
| 5 | us | North America |
| 6 | us | North America |

case_when() with value if all tests fail (use a variable!)

```
ufo <- ufo %>% mutate(  
  region = case_when(  
    country %in% "gb" ~ "Great Brit.",  
    TRUE ~ region  
  ))  
ufo %>% select(country, region) %>% head()
```

```
# A tibble: 6 x 2  
  country region  
  <chr>    <chr>  
1 us      North America  
2 <NA>    Other  
3 gb      Great Brit.  
4 us      North America  
5 us      North America  
6 us      North America
```

Ordering the rows of a data.frame: dplyr

The arrange function can reorder rows By default, arrange orders in ascending order:

```
ufo %>% arrange(duration_s)
```

```
# A tibble: 88,875 x 12
```

| | datetime | City | state | country | shape | durat~1 | durat~2 | co |
|----|-------------|-------|-------|---------|-------|---------|---------|-------|
| | <chr> | <chr> | <chr> | <chr> | <chr> | <dbl> | <chr> | <chr> |
| 1 | 10/10/1995~ | puer~ | pr | <NA> | <NA> | 0 | <NA> | Wo |
| 2 | 10/10/1999~ | ashl~ | mo | us | light | 0 | two se~ | We |
| 3 | 10/10/2002~ | baha~ | <NA> | <NA> | egg | 0 | <NA> | we |
| 4 | 10/10/2002~ | burn~ | <NA> | au | cross | 0 | 12 | th |
| 5 | 10/10/2005~ | edge~ | fl | us | <NA> | 0 | 300 | or |
| 6 | 10/10/2005~ | fran~ | in | us | disk | 0 | ? | tw |
| 7 | 10/10/2006~ | knik | ak | us | tria~ | 0 | 5 | SI |
| 8 | 10/10/2007~ | bake~ | ca | us | circ~ | 0 | had a ~ | UF |
| 9 | 10/10/2008~ | amar~ | tx | us | flash | 0 | <NA> | we |
| 10 | 10/10/2009~ | gree~ | <NA> | <NA> | rect~ | 0 | <NA> | Fe |

```
# ... with 88,865 more rows, 2 more variables: longitude <dbl>
```

Ordering the rows of a data.frame: dplyr

Use the desc to arrange the rows in descending order:

```
ufo %>% arrange(desc(duration_s))
```

```
# A tibble: 88,875 x 12
```

| | datetime | City | state | country | shape | durat~1 | durat~2 | co |
|----|-----------|----------|-------|---------|-------|---------|---------|-------|
| | <chr> | <chr> | <chr> | <chr> | <chr> | <dbl> | <chr> | <chr> |
| 1 | 10/1/1983 | ~ birm~ | <NA> | gb | sphe~ | 9.78e7 | 31 yea~ | F |
| 2 | 6/3/2010 | 2~ otta~ | on | ca | other | 8.28e7 | 23000h~ | (|
| 3 | 9/15/1991 | ~ gree~ | ar | us | light | 6.63e7 | 21 yea~ | On |
| 4 | 4/2/1983 | 2~ dont~ | <NA> | <NA> | <NA> | 5.26e7 | 2 mont~ | H |
| 5 | 8/10/2012 | ~ finl~ | wa | us | light | 5.26e7 | 2 mont~ | Th |
| 6 | 8/24/2002 | ~ engl~ | fl | us | light | 5.26e7 | 2 mont~ | br |
| 7 | 6/30/1969 | ~ some~ | <NA> | gb | cone | 2.52e7 | 8 years | F |
| 8 | 10/7/2013 | ~ okla~ | ok | <NA> | circ~ | 1.05e7 | 4 mont~ | Br |
| 9 | 3/1/1994 | 0~ meni~ | ca | us | unkn~ | 1.05e7 | 4 mont~ | Su |
| 10 | 8/3/2008 | 2~ virg~ | va | us | fire~ | 1.05e7 | 4 mont~ | th |

```
# ... with 88,865 more rows, 2 more variables: longitude <dbl>, latitude <dbl>  
# and abbreviated variable names 1: duration_s, 2: `duration_s`
```

Ordering the rows of a data.frame: dplyr

It is a bit more straightforward to mix increasing and decreasing orderings:

```
ufo %>% arrange(country, desc(duration_s))
```

```
# A tibble: 88,875 x 12
```

| | datetime | City | state | country | shape | durat~1 | durat~2 | co |
|----|-------------|-------|-------|---------|-------|---------|---------|-------|
| | <chr> | <chr> | <chr> | <chr> | <chr> | <dbl> | <chr> | <chr> |
| 1 | 11/12/2013~ | moun~ | <NA> | au | sphe~ | 1209600 | 2 weeks | Or |
| 2 | 5/12/2004 ~ | sydn~ | <NA> | au | light | 345600 | 4 days+ | In |
| 3 | 4/18/2004 ~ | sydn~ | <NA> | au | light | 86400 | day | It |
| 4 | 4/15/1983 ~ | bris~ | <NA> | au | chan~ | 37800 | 1 1/2 ~ | A |
| 5 | 4/18/1996 ~ | bris~ | <NA> | au | <NA> | 18000 | 5 hou~ | F |
| 6 | 6/9/2005 2~ | melb~ | <NA> | au | circ~ | 18000 | 5 hour~ | UF |
| 7 | 11/6/2009 ~ | pert~ | <NA> | au | light | 14400 | 4hrs | Un |
| 8 | 3/15/2004 ~ | adel~ | <NA> | au | form~ | 10800 | 1-3 hrs | iv |
| 9 | 3/2/2014 2~ | pert~ | <NA> | au | light | 10800 | 2-3 ho~ | Co |
| 10 | 6/20/2001 ~ | canb~ | <NA> | au | tear~ | 10800 | 3 hrs | 8 |

```
# ... with 88,865 more rows, 2 more variables: longitude <dbl>
```

Lab

Link to Lab