

Data Summarization

Data Wrangling in R

Data Summarization

- Basic statistical summarization
 - `mean(x)`: takes the mean of x
 - `sd(x)`: takes the standard deviation of x
 - `median(x)`: takes the median of x
 - `quantile(x)`: displays sample quantiles of x. Default is min, IQR, max
 - `range(x)`: displays the range. Same as `c(min(x), max(x))`
 - `sum(x)`: sum of x
 - `max(x)`: maximum value in x
 - `min(x)`: minimum value in x
- **all have the `na.rm` = argument for missing data**

Statistical summarization

These functions work on **vectors**:

```
x <- c(1, 5, 7, 4, 2, 8)
mean(x)
```

```
[1] 4.5
```

```
mean(x, na.rm = TRUE) # Remove NAs if needed
```

```
[1] 4.5
```

Statistical summarization

Summarization on a `data.frame`/tibble:

```
mtcars %>% pull(hp) %>% mean() # alt: pull(mtcars, hp) %>% mean()
```

```
[1] 146.6875
```

```
mean(mtcars$hp)
```

```
[1] 146.6875
```

Youth Tobacco Survey

Here we will be using the Youth Tobacco Survey data:

```
yts <-  
  read_csv("http://jhudatascience.org/intro_to_r/data/Youth_Tobacco_Survey_YTS_Data.csv")  
head(yts)
```

```
# A tibble: 6 × 31  
  YEAR LocationAbbr LocationDesc TopicType      TopicDesc MeasureDesc DataSource  
  <dbl> <chr>      <chr>      <chr>      <chr>      <chr>      <chr>  
1  2015 AZ          Arizona    Tobacco Use ... Cessatio... Percent of... YTS  
2  2015 AZ          Arizona    Tobacco Use ... Cessatio... Percent of... YTS  
3  2015 AZ          Arizona    Tobacco Use ... Cessatio... Percent of... YTS  
4  2015 AZ          Arizona    Tobacco Use ... Cessatio... Quit Attem... YTS  
5  2015 AZ          Arizona    Tobacco Use ... Cessatio... Quit Attem... YTS  
6  2015 AZ          Arizona    Tobacco Use ... Cessatio... Quit Attem... YTS  
# ... with 24 more variables: Response <chr>, Data_Value_Unit <chr>,  
#   Data_Value_Type <chr>, Data_Value <dbl>, Data_Value_Footnote_Symbol <chr>,  
#   Data_Value_Footnote <chr>, Data_Value_Std_Err <dbl>,  
#   Low_Confidence_Limit <dbl>, High_Confidence_Limit <dbl>, Sample_Size <dbl>,  
#   Gender <chr>, Race <chr>, Age <chr>, Education <chr>, GeoLocation <chr>,  
#   TopicTypeId <chr>, TopicId <chr>, MeasureId <chr>, StratificationID1 <chr>,  
#   StratificationID2 <chr>, StratificationID3 <chr>, ...
```

Column to vector

Let's work with one column as a vector using `pull()`.

```
locations <- yts %>% pull(LocationDesc)
locations
```

[1]	"Arizona"	"Arizona"
[3]	"Arizona"	"Arizona"
[5]	"Arizona"	"Arizona"
[7]	"Arizona"	"Arizona"
[9]	"Arizona"	"Arizona"
[11]	"Arizona"	"Arizona"
[13]	"Arizona"	"Arizona"
[15]	"Arizona"	"Arizona"
[17]	"Arizona"	"Arizona"
[19]	"Arizona"	"Arizona"
[21]	"Arizona"	"Arizona"
[23]	"Arizona"	"Arizona"
[25]	"Connecticut"	"Connecticut"
[27]	"Connecticut"	"Connecticut"
[29]	"Connecticut"	"Connecticut"
[31]	"Connecticut"	"Connecticut"
[33]	"Connecticut"	"Connecticut"
[35]	"Connecticut"	"Connecticut"
[37]	"Connecticut"	"Connecticut"
[39]	"Connecticut"	"Connecticut"
[41]	"Connecticut"	"Connecticut"
[43]	"Connecticut"	"Connecticut"
[45]	"Connecticut"	"Connecticut"

Length and unique

`unique(x)` will return the unique elements of `x`

```
unique(locations)
```

```
[1] "Arizona"           "Connecticut"
[3] "Georgia"           "Hawaii"
[5] "Illinois"          "Louisiana"
[7] "Mississippi"       "Utah"
[9] "Missouri"          "National (States and DC) "
[11] "Nebraska"          "New Jersey"
[13] "North Carolina"    "North Dakota"
[15] "Pennsylvania"      "South Carolina"
[17] "West Virginia"     "Alabama"
[19] "Delaware"          "Minnesota"
[21] "Guam"              "Ohio"
[23] "Indiana"           "Kansas"
[25] "Oklahoma"          "Wisconsin"
[27] "Michigan"          "New Hampshire"
[29] "Arkansas"          "Kentucky"
[31] "Iowa"              "South Dakota"
[33] "Virginia"          "Puerto Rico"
[35] "Rhode Island"      "New Mexico"
[37] "Tennessee"         "Vermont"
[39] "Virgin Islands"    "California"
[41] "Idaho"             "Florida"
[43] "Maryland"          "Massachusetts"
[45] "New York"          "Maine"
```

Length and unique

`length` will tell you the length of a vector. Combined with `unique`, tells you the number of unique elements:

```
length(unique(locations))
```

```
[1] 50
```


table and dplyr: count

table(x) will return a frequency table of unique elements of x

```
table(locations)
```

locations

Alabama	Arizona	Arkansas
378	240	210
California	Colorado	Connecticut
96	48	384
Delaware	District of Columbia	Florida
312	48	96
Georgia	Guam	Hawaii
282	48	270
Idaho	Illinois	Indiana
48	282	264
Iowa	Kansas	Kentucky
276	186	255
Louisiana	Maine	Maryland
240	48	96
Massachusetts	Michigan	Minnesota
48	138	141
Mississippi	Missouri	National (States and DC)
567	294	26
Nebraska	New Hampshire	New Jersey
234	180	387
New Mexico	New York	North Carolina
24	90	366

table and dplyr: count

Use `count` directly on a `data.frame` and column without needing to use `pull()`.

```
yts %>% count(LocationDesc)
```

```
# A tibble: 50 × 2
```

	LocationDesc	n
	<chr>	<int>
1	Alabama	378
2	Arizona	240
3	Arkansas	210
4	California	96
5	Colorado	48
6	Connecticut	384
7	Delaware	312
8	District of Columbia	48
9	Florida	96
10	Georgia	282
# ... with 40 more rows		

table and dplyr: count

Multiple columns listed further subdivides the count.

```
yts %>% count(LocationDesc, TopicDesc)
```

```
# A tibble: 146 × 3
```

	LocationDesc	TopicDesc	n
	<chr>	<chr>	<int>
1	Alabama	Cessation (Youth)	90
2	Alabama	Cigarette Use (Youth)	144
3	Alabama	Smokeless Tobacco Use (Youth)	144
4	Arizona	Cessation (Youth)	60
5	Arizona	Cigarette Use (Youth)	99
6	Arizona	Smokeless Tobacco Use (Youth)	81
7	Arkansas	Cessation (Youth)	42
8	Arkansas	Cigarette Use (Youth)	78
9	Arkansas	Smokeless Tobacco Use (Youth)	90
10	California	Cessation (Youth)	24

```
# ... with 136 more rows
```

table and dplyr: count

Multiple columns listed further subdivides the count.

```
yts %>% count(LocationDesc, TopicDesc)
```

```
# A tibble: 146 × 3
```

	LocationDesc	TopicDesc	n
	<chr>	<chr>	<int>
1	Alabama	Cessation (Youth)	90
2	Alabama	Cigarette Use (Youth)	144
3	Alabama	Smokeless Tobacco Use (Youth)	144
4	Arizona	Cessation (Youth)	60
5	Arizona	Cigarette Use (Youth)	99
6	Arizona	Smokeless Tobacco Use (Youth)	81
7	Arkansas	Cessation (Youth)	42
8	Arkansas	Cigarette Use (Youth)	78
9	Arkansas	Smokeless Tobacco Use (Youth)	90
10	California	Cessation (Youth)	24

```
# ... with 136 more rows
```

table and dplyr: count

Note: `count()` includes NAs but `table()` does not

Grouping

Perform Operations By Groups: dplyr

`group_by` allows you group the data set by variables/columns you specify:

```
# Regular data  
yts
```

```
# A tibble: 9,794 × 31  
  YEAR LocationAbbr LocationDesc TopicType TopicDesc MeasureDesc DataSource  
  <dbl> <chr>      <chr>      <chr>      <chr>      <chr>      <chr>  
1  2015 AZ          Arizona    Tobacco Use... Cessatio... Percent of... YTS  
2  2015 AZ          Arizona    Tobacco Use... Cessatio... Percent of... YTS  
3  2015 AZ          Arizona    Tobacco Use... Cessatio... Percent of... YTS  
4  2015 AZ          Arizona    Tobacco Use... Cessatio... Quit Attem... YTS  
5  2015 AZ          Arizona    Tobacco Use... Cessatio... Quit Attem... YTS  
6  2015 AZ          Arizona    Tobacco Use... Cessatio... Quit Attem... YTS  
7  2015 AZ          Arizona    Tobacco Use... Cigarette... Smoking St... YTS  
8  2015 AZ          Arizona    Tobacco Use... Cigarette... Smoking St... YTS  
9  2015 AZ          Arizona    Tobacco Use... Cigarette... Smoking St... YTS  
10 2015 AZ          Arizona    Tobacco Use... Cigarette... Smoking St... YTS  
# ... with 9,784 more rows, and 24 more variables: Response <chr>,  
# Data_Value_Unit <chr>, Data_Value_Type <chr>, Data_Value <dbl>,  
# Data_Value_Footnote_Symbol <chr>, Data_Value_Footnote <chr>,  
# Data_Value_Std_Err <dbl>, Low_Confidence_Limit <dbl>,  
# High_Confidence_Limit <dbl>, Sample_Size <dbl>, Gender <chr>, Race <chr>,  
# Age <chr>, Education <chr>, GeoLocation <chr>, TopicTypeId <chr>,  
# TopicId <chr>, MeasureId <chr>, StratificationID1 <chr>, ...
```

Perform Operations By Groups: dplyr

`group_by` allows you group the data set by variables/columns you specify:

```
yts_grouped <- yts %>% group_by(Response)
yts_grouped
```

```
# A tibble: 9,794 × 31
# Groups:   Response [4]
  YEAR LocationAbbr LocationDesc TopicType TopicDesc MeasureDesc DataSource
  <dbl> <chr>         <chr>         <chr>         <chr>         <chr>         <chr>
1  2015 AZ          Arizona      Tobacco Use... Cessatio... Percent of... YTS
2  2015 AZ          Arizona      Tobacco Use... Cessatio... Percent of... YTS
3  2015 AZ          Arizona      Tobacco Use... Cessatio... Percent of... YTS
4  2015 AZ          Arizona      Tobacco Use... Cessatio... Quit Attem... YTS
5  2015 AZ          Arizona      Tobacco Use... Cessatio... Quit Attem... YTS
6  2015 AZ          Arizona      Tobacco Use... Cessatio... Quit Attem... YTS
7  2015 AZ          Arizona      Tobacco Use... Cigarette... Smoking St... YTS
8  2015 AZ          Arizona      Tobacco Use... Cigarette... Smoking St... YTS
9  2015 AZ          Arizona      Tobacco Use... Cigarette... Smoking St... YTS
10 2015 AZ          Arizona      Tobacco Use... Cigarette... Smoking St... YTS
# ... with 9,784 more rows, and 24 more variables: Response <chr>,
# Data_Value_Unit <chr>, Data_Value_Type <chr>, Data_Value <dbl>,
# Data_Value_Footnote_Symbol <chr>, Data_Value_Footnote <chr>,
# Data_Value_Std_Err <dbl>, Low_Confidence_Limit <dbl>,
# High_Confidence_Limit <dbl>, Sample_Size <dbl>, Gender <chr>, Race <chr>,
# Age <chr>, Education <chr>, GeoLocation <chr>, TopicTypeId <chr>,
# TopicId <chr>, MeasureId <chr>, StratificationID1 <chr>, ...
```


Summarize the data: `dplyr` `summarize()` function

`summarize` is a helpful function to use after `group_by()`. It creates a summary table of a column you're interested in.

```
yts %>%  
  summarize(mean_value = mean(Data_Value, na.rm = TRUE))
```

```
# A tibble: 1 × 1  
  mean_value  
    <dbl>  
1      21.0
```

Summarize the grouped data

It's grouped! Grouping doesn't change the data in any way, but how **functions operate on it**. Now we can summarize `Data_Value` (percent of respondents) by group:

```
yts_grouped %>%  
  summarize(mean_value = mean(Data_Value, na.rm = TRUE))
```

```
# A tibble: 4 × 2  
  Response mean_value  
  <chr>      <dbl>  
1 Current      9.68  
2 Ever        26.1  
3 Frequent     3.48  
4 <NA>        53.5
```

Use the **pipe** to string these together!

Pipe `yts` into `group_by`, then pipe that into `summarize`:

```
yts %>%  
  group_by(Response) %>%  
  summarize(mean_value = mean(Data_Value, na.rm = TRUE),  
            max_value = max(Data_Value, na.rm = TRUE))
```

```
# A tibble: 4 × 3  
  Response mean_value max_value  
  <chr>      <dbl>      <dbl>  
1 Current      9.68      40.6  
2 Ever        26.1       98  
3 Frequent     3.48      23.9  
4 <NA>        53.5      81.9
```

group_by with mutate - just add data

We can also use `mutate` to calculate the mean value for each year and add it as a column:

```
yts %>%  
  group_by(YEAR) %>%  
  mutate(year_avg = mean(Data_Value, na.rm = TRUE)) %>%  
  select(LocationDesc, Data_Value, year_avg)
```

```
# A tibble: 9,794 × 4  
# Groups:   YEAR [17]  
   YEAR LocationDesc Data_Value year_avg  
   <dbl> <chr>         <dbl>    <dbl>  
1  2015 Arizona          NA      15.2  
2  2015 Arizona          NA      15.2  
3  2015 Arizona          NA      15.2  
4  2015 Arizona          NA      15.2  
5  2015 Arizona          NA      15.2  
6  2015 Arizona          NA      15.2  
7  2015 Arizona         3.2      15.2  
8  2015 Arizona         3.2      15.2  
9  2015 Arizona         3.1      15.2  
10 2015 Arizona        12.5      15.2  
# ... with 9,784 more rows
```

Counting

There are other functions, such as `n()` count the number of observations.

```
yts %>%  
  group_by(YEAR) %>%  
  summarize(n = n(),  
            mean = mean(Data_Value, na.rm = TRUE))
```

```
# A tibble: 17 × 3  
  YEAR      n mean  
  <dbl> <int> <dbl>  
1  1999   372  26.1  
2  2000  1224  26.7  
3  2001   426  23.4  
4  2002  1016  25.2  
5  2003   498  21.3  
6  2004   611  20.7  
7  2005   636  21.8  
8  2006   518  21.8  
9  2007   516  20.0  
10 2008   483  18.2  
11 2009   686  18.3  
12 2010   447  17.8  
13 2011   521  17.8  
14 2012   244  15.5  
15 2013   685  16.7  
16 2014   334  15.7  
17 2015   577  15.2
```

Iterative summaries

Iterative summaries: `dplyr` `summarize()` and `across()` functions

Use the `across` function with `summarize()` to summarize across multiple columns of your data.

```
# General format - Not the code!
```

```
across({ columns to go across }, ~ { summarization_function(.x, na.rm = ..) })
```

```
yts %>%
```

```
  group_by(YEAR) %>%
```

```
  summarize(across( c(Data_Value, Data_Value_Std_Err, Sample_Size),  
                    ~ mean(.x, na.rm = TRUE)))
```

```
# A tibble: 17 × 4
```

	YEAR	Data_Value	Data_Value_Std_Err	Sample_Size
	<dbl>	<dbl>	<dbl>	<dbl>
1	1999	26.1	1.98	1591.
2	2000	26.7	2.03	1743.
3	2001	23.4	1.79	2060.
4	2002	25.2	1.81	2653.
5	2003	21.3	1.92	2325.
6	2004	20.7	1.84	1246.
7	2005	21.8	2.17	1017.
8	2006	21.8	2.15	1191.
9	2007	20.0	1.96	1093.
10	2008	18.2	1.73	1203.
11	2009	18.3	1.90	1033.
12	2010	17.8	1.71	1202.
13	2011	17.8	1.84	1274.

Iterative summaries: `dplyr::summarize()` and `across()` functions

Another example using select helpers (??`tidyr::tidy_select`):

```
yts %>%
  summarize(across( starts_with("Data"), ~ range(.x, na.rm = TRUE)))
```

A tibble: 2 × 7

	DataSource	Data_Value_Unit	Data_Value_Type	Data_Value	Data_Value_Footnote_Sym...
	<chr>	<chr>	<chr>	<dbl>	<chr>
1	NYTS	%	Percentage	0	*
2	YTS	%	Percentage	98	*

... with 2 more variables: Data_Value_Footnote <chr>, Data_Value_Std_Err <dbl>

Data Summarization on data frames

- Basic statistical summarization for numeric data
 - `rowMeans(x)`: takes the means of each row of x
 - `colMeans(x)`: takes the means of each column of x
 - `rowSums(x)`: takes the sum of each row of x
 - `colSums(x)`: takes the sum of each column of x
 - `summary(x)`: for data frames, displays the quantile information

summary() Function

Using `summary()` can give you rough snapshots of each numeric column (character columns are skipped):

```
summary(yts)
```

YEAR	LocationAbbr	LocationDesc	TopicType
Min. :1999	Length:9794	Length:9794	Length:9794
1st Qu.:2002	Class :character	Class :character	Class :character
Median :2006	Mode :character	Mode :character	Mode :character
Mean :2006			
3rd Qu.:2010			
Max. :2015			

TopicDesc	MeasureDesc	DataSource	Response
Length:9794	Length:9794	Length:9794	Length:9794
Class :character	Class :character	Class :character	Class :character
Mode :character	Mode :character	Mode :character	Mode :character

Data_Value_Unit	Data_Value_Type	Data_Value
Length:9794	Length:9794	Min. : 0.00
Class :character	Class :character	1st Qu.: 3.20
Mode :character	Mode :character	Median :11.30
		Mean :20.97
		3rd Qu.:39.10

Summary

- summary stats (`mean()`) work with `pull()`
- `count(x)`: what unique values do you have?
 - `pull() + table()`
 - `unique()` combined with `length()`
- `group_by()`: changes all subsequent functions
 - combine with `summarize()` to get statistics per group
 - combine with `across()` to programmatically select columns
- `summary(x)`: quantile information