

Data Summarization

Data Wrangling in R

Data Summarization

- Basic statistical summarization
 - `mean(x)`: takes the mean of x
 - `sd(x)`: takes the standard deviation of x
 - `median(x)`: takes the median of x
 - `quantile(x)`: displays sample quantiles of x. Default is min, IQR, max
 - `range(x)`: displays the range. Same as `c(min(x), max(x))`
 - `sum(x)`: sum of x
 - `max(x)`: maximum value in x
 - `min(x)`: minimum value in x
- **all have the `na.rm` = argument for missing data**

Statistical summarization

These functions work on **vectors**:

```
x <- c(1, 5, 7, 4, 2, 8)
mean(x)
```

```
[1] 4.5
```

```
mean(x, na.rm = TRUE) # Remove NAs if needed
```

```
[1] 4.5
```

Statistical summarization

Summarization on a `data.frame`/tibble:

```
mtcars %>% pull(hp) %>% mean() # alt: pull(mtcars, hp) %>% mean()
```

```
[1] 146.6875
```

```
mean(mtcars$hp)
```

```
[1] 146.6875
```

Youth Tobacco Survey

Let's use the Youth Tobacco Survey data again:

```
yts <-  
  read_csv("https://sisbid.github.io/Data-Wrangling/data/Youth_Tobacco_Survey_YTS_Data.csv")  
head(yts)  
  
# A tibble: 6 × 31  
  YEAR LocationAbbr LocationDesc TopicType TopicDesc MeasureDesc DataSource Response Data_Value_Unit  
  <dbl> <chr>      <chr>      <chr>      <chr>      <chr>      <chr>      <chr>      <chr>  
1  2015 AZ          Arizona Tobacco Use - Sur... Cessatio... Percent of... YTS      <NA>      %  
2  2015 AZ          Arizona Tobacco Use - Sur... Cessatio... Percent of... YTS      <NA>      %  
3  2015 AZ          Arizona Tobacco Use - Sur... Cessatio... Percent of... YTS      <NA>      %  
4  2015 AZ          Arizona Tobacco Use - Sur... Cessatio... Quit Attem... YTS      <NA>      %  
5  2015 AZ          Arizona Tobacco Use - Sur... Cessatio... Quit Attem... YTS      <NA>      %  
6  2015 AZ          Arizona Tobacco Use - Sur... Cessatio... Quit Attem... YTS      <NA>      %  
# i 22 more variables: Data_Value_Type <chr>, Data_Value <dbl>, Data_Value_Footnote_Symbol <chr>,  
# Data_Value_Footnote <chr>, Data_Value_Std_Err <dbl>, Low_Confidence_Limit <dbl>,  
# High_Confidence_Limit <dbl>, Sample_Size <dbl>, Gender <chr>, Race <chr>, Age <chr>, Education <chr>,  
# GeoLocation <chr>, TopicTypeId <chr>, TopicId <chr>, MeasureId <chr>, StratificationID1 <chr>,  
# StratificationID2 <chr>, StratificationID3 <chr>, StratificationID4 <chr>, SubMeasureID <chr>,  
# DisplayOrder <dbl>
```

Column to vector

Let's work with one column as a vector using `pull()`.

```
locations <- yts %>% pull(LocationDesc)
locations
```

[1]	"Arizona"	"Arizona"	"Arizona"
[4]	"Arizona"	"Arizona"	"Arizona"
[7]	"Arizona"	"Arizona"	"Arizona"
[10]	"Arizona"	"Arizona"	"Arizona"
[13]	"Arizona"	"Arizona"	"Arizona"
[16]	"Arizona"	"Arizona"	"Arizona"
[19]	"Arizona"	"Arizona"	"Arizona"
[22]	"Arizona"	"Arizona"	"Arizona"
[25]	"Connecticut"	"Connecticut"	"Connecticut"
[28]	"Connecticut"	"Connecticut"	"Connecticut"
[31]	"Connecticut"	"Connecticut"	"Connecticut"
[34]	"Connecticut"	"Connecticut"	"Connecticut"
[37]	"Connecticut"	"Connecticut"	"Connecticut"
[40]	"Connecticut"	"Connecticut"	"Connecticut"
[43]	"Connecticut"	"Connecticut"	"Connecticut"
[46]	"Connecticut"	"Connecticut"	"Connecticut"
[49]	"Connecticut"	"Connecticut"	"Connecticut"
[52]	"Connecticut"	"Connecticut"	"Connecticut"
[55]	"Connecticut"	"Connecticut"	"Connecticut"
[58]	"Connecticut"	"Connecticut"	"Connecticut"
[61]	"Connecticut"	"Connecticut"	"Connecticut"
[64]	"Connecticut"	"Connecticut"	"Connecticut"
[67]	"Connecticut"	"Connecticut"	"Connecticut"

Length and unique

`unique(x)` will return the unique elements of `x`

```
unique(locations)
```

```
[1] "Arizona"           "Connecticut"      "Georgia"
[4] "Hawaii"            "Illinois"         "Louisiana"
[7] "Mississippi"       "Utah"             "Missouri"
[10] "National (States and DC)" "Nebraska"         "New Jersey"
[13] "North Carolina"    "North Dakota"     "Pennsylvania"
[16] "South Carolina"    "West Virginia"    "Alabama"
[19] "Delaware"          "Minnesota"        "Guam"
[22] "Ohio"              "Indiana"          "Kansas"
[25] "Oklahoma"          "Wisconsin"        "Michigan"
[28] "New Hampshire"     "Arkansas"         "Kentucky"
[31] "Iowa"              "South Dakota"     "Virginia"
[34] "Puerto Rico"      "Rhode Island"     "New Mexico"
[37] "Tennessee"        "Vermont"          "Virgin Islands"
[40] "California"        "Idaho"            "Florida"
[43] "Maryland"          "Massachusetts"    "New York"
[46] "Maine"             "Colorado"         "District of Columbia"
[49] "Texas"             "Wyoming"
```

Length and unique

`length` will tell you the length of a vector. Combined with `unique`, tells you the number of unique elements:

```
length(unique(locations))
```

```
[1] 50
```


Counting NAs

```
use sum(is.na()):
```

```
sum(is.na(locations))
```

```
[1] 0
```

dplyr: count

Use `count` directly on a `data.frame` and `column`: count the number of rows in each group.

```
yts %>% count(LocationDesc)
```

```
# A tibble: 50 × 2
```

	LocationDesc	n
	<chr>	<int>
1	Alabama	378
2	Arizona	240
3	Arkansas	210
4	California	96
5	Colorado	48
6	Connecticut	384
7	Delaware	312
8	District of Columbia	48
9	Florida	96
10	Georgia	282

```
# i 40 more rows
```

dplyr::count

Multiple columns listed further subdivides the count.

```
yts %>% count(LocationDesc, TopicDesc)
```

```
# A tibble: 146 × 3
  LocationDesc TopicDesc      n
  <chr>        <chr>    <int>
1 Alabama      Cessation (Youth)      90
2 Alabama      Cigarette Use (Youth) 144
3 Alabama      Smokeless Tobacco Use (Youth) 144
4 Arizona      Cessation (Youth)      60
5 Arizona      Cigarette Use (Youth)   99
6 Arizona      Smokeless Tobacco Use (Youth) 81
7 Arkansas     Cessation (Youth)      42
8 Arkansas     Cigarette Use (Youth)   78
9 Arkansas     Smokeless Tobacco Use (Youth) 90
10 California  Cessation (Youth)      24
# i 136 more rows
```

dplyr: count

Option to sort the results with `sort = TRUE`

```
yts %>% count(LocationDesc, sort = TRUE)
```

```
# A tibble: 50 × 2
```

	LocationDesc	n
	<chr>	<int>
1	Mississippi	567
2	New Jersey	387
3	Connecticut	384
4	Alabama	378
5	North Carolina	366
6	Wisconsin	360
7	West Virginia	336
8	North Dakota	330
9	Pennsylvania	330
10	Oklahoma	318

```
# i 40 more rows
```

dplyr: count

Instead of counting the number of rows in each group, `wt` computes `sum(wt)` for each group.

```
# Add up "Data_Value" for each LocationDesc category
yts %>% count(LocationDesc, wt = Data_Value)
```

```
# A tibble: 50 × 2
  LocationDesc      n
  <chr>          <dbl>
1 Alabama      9220.
2 Arizona      3937.
3 Arkansas      5443.
4 California    2059.
5 Colorado      1136.
6 Connecticut    5838.
7 Delaware      5886
8 District of Columbia  853.
9 Florida       2786.
10 Georgia      5625.
# i 40 more rows
```

Grouping

Perform Operations By Groups: dplyr

`group_by` allows you group the data set by variables/columns you specify:

```
# Regular data  
yts
```

```
# A tibble: 9,794 × 31
```

	YEAR	LocationAbbr	LocationDesc	TopicType	TopicDesc	MeasureDesc	DataSource	Response
	<dbl>	<chr>	<chr>	<chr>	<chr>	<chr>	<chr>	<chr>
1	2015	AZ	Arizona	Tobacco Use - Su...	Cessatio...	Percent of...	YTS	<NA>
2	2015	AZ	Arizona	Tobacco Use - Su...	Cessatio...	Percent of...	YTS	<NA>
3	2015	AZ	Arizona	Tobacco Use - Su...	Cessatio...	Percent of...	YTS	<NA>
4	2015	AZ	Arizona	Tobacco Use - Su...	Cessatio...	Quit Attem..	YTS	<NA>
5	2015	AZ	Arizona	Tobacco Use - Su...	Cessatio...	Quit Attem..	YTS	<NA>
6	2015	AZ	Arizona	Tobacco Use - Su...	Cessatio...	Quit Attem..	YTS	<NA>
7	2015	AZ	Arizona	Tobacco Use - Su...	Cigarett...	Smoking St...	YTS	Current
8	2015	AZ	Arizona	Tobacco Use - Su...	Cigarett...	Smoking St...	YTS	Current
9	2015	AZ	Arizona	Tobacco Use - Su...	Cigarett...	Smoking St...	YTS	Current
10	2015	AZ	Arizona	Tobacco Use - Su...	Cigarett...	Smoking St...	YTS	Ever

```
# i 9,784 more rows
```

```
# i 22 more variables: Data_Value_Type <chr>, Data_Value <dbl>, Data_Value_Footnote_Symbol <chr>  
# Data_Value_Footnote <chr>, Data_Value_Std_Err <dbl>, Low_Confidence_Limit <dbl>,  
# High_Confidence_Limit <dbl>, Sample_Size <dbl>, Gender <chr>, Race <chr>, Age <chr>, Educat  
# GeoLocation <chr>, TopicTypeId <chr>, TopicId <chr>, MeasureId <chr>, StratificationID1 <chr>  
# StratificationID2 <chr>, StratificationID3 <chr>, StratificationID4 <chr>, SubMeasureID <chr>  
# DisplayOrder <dbl>
```

Perform Operations By Groups: dplyr

`group_by` allows you group the data set by variables/columns you specify:

```
yts_grouped <- yts %>% group_by(Response)
yts_grouped
```

```
# A tibble: 9,794 × 31
```

```
# Groups:   Response [4]
```

	YEAR	LocationAbbr	LocationDesc	TopicType	TopicDesc	MeasureDesc	DataSource	Response
	<dbl>	<chr>	<chr>	<chr>	<chr>	<chr>	<chr>	<chr>
1	2015	AZ	Arizona	Tobacco Use - Su...	Cessatio...	Percent of...	YTS	<NA>
2	2015	AZ	Arizona	Tobacco Use - Su...	Cessatio...	Percent of...	YTS	<NA>
3	2015	AZ	Arizona	Tobacco Use - Su...	Cessatio...	Percent of...	YTS	<NA>
4	2015	AZ	Arizona	Tobacco Use - Su...	Cessatio...	Quit Attem..	YTS	<NA>
5	2015	AZ	Arizona	Tobacco Use - Su...	Cessatio...	Quit Attem..	YTS	<NA>
6	2015	AZ	Arizona	Tobacco Use - Su...	Cessatio...	Quit Attem..	YTS	<NA>
7	2015	AZ	Arizona	Tobacco Use - Su...	Cigarett...	Smoking St...	YTS	Current
8	2015	AZ	Arizona	Tobacco Use - Su...	Cigarett...	Smoking St...	YTS	Current
9	2015	AZ	Arizona	Tobacco Use - Su...	Cigarett...	Smoking St...	YTS	Current
10	2015	AZ	Arizona	Tobacco Use - Su...	Cigarett...	Smoking St...	YTS	Ever

```
# i 9,784 more rows
```

```
# i 22 more variables: Data_Value_Type <chr>, Data_Value <dbl>, Data_Value_Footnote_Symbol <chr>,
# Data_Value_Footnote <chr>, Data_Value_Std_Err <dbl>, Low_Confidence_Limit <dbl>,
# High_Confidence_Limit <dbl>, Sample_Size <dbl>, Gender <chr>, Race <chr>, Age <chr>, Educat
# GeoLocation <chr>, TopicTypeId <chr>, TopicId <chr>, MeasureId <chr>, StratificationID1 <chr>,
# StratificationID2 <chr>, StratificationID3 <chr>, StratificationID4 <chr>, SubMeasureID <chr>,
# DisplayOrder <dbl>
```


Summarize the data: `dplyr summarize()` function

`summarize` is a helpful function to use after `group_by()`. It creates a summary table of a column you're interested in.

```
yts %>%  
  summarize(mean_value = mean(Data_Value, na.rm = TRUE))
```

```
# A tibble: 1 × 1  
  mean_value  
    <dbl>  
1      21.0
```

Summarize the grouped data

It's grouped! Grouping doesn't change the data in any way, but how **functions operate on it**. Now we can summarize `Data_Value` (percent of respondents) by group:

```
yts_grouped %>%  
  summarize(mean_value = mean(Data_Value, na.rm = TRUE))
```

```
# A tibble: 4 × 2  
  Response mean_value  
  <chr>         <dbl>  
1 Current      9.68  
2 Ever        26.1  
3 Frequent     3.48  
4 <NA>        53.5
```

Use the **pipe** to string these together!

Pipe `yts` into `group_by`, then pipe that into `summarize`:

```
yts %>%  
  group_by(Response) %>%  
  summarize(mean_value = mean(Data_Value, na.rm = TRUE),  
            max_value = max(Data_Value, na.rm = TRUE))
```

```
# A tibble: 4 × 3  
  Response mean_value max_value  
  <chr>      <dbl>      <dbl>  
1 Current      9.68      40.6  
2 Ever        26.1       98  
3 Frequent     3.48      23.9  
4 <NA>        53.5      81.9
```

group_by with **mutate** - Useful for comparisons

Use `group_by` to calculate the mean value for each year. We can use `mutate` to add it as a column.

```
yts_year <- yts %>%  
  group_by(YEAR) %>%  
  mutate(year_avg = mean(Data_Value, na.rm = TRUE)) %>%  
  select(LocationDesc, Data_Value, year_avg)
```

group_by with mutate - Useful for comparisons

Create a “difference” variable:

```
yts_year %>% mutate(Diff = Data_Value - year_avg)
```

```
# A tibble: 9,794 × 5
```

```
# Groups:   YEAR [17]
```

	YEAR	LocationDesc	Data_Value	year_avg	Diff
	<dbl>	<chr>	<dbl>	<dbl>	<dbl>
1	2015	Arizona	NA	15.2	NA
2	2015	Arizona	NA	15.2	NA
3	2015	Arizona	NA	15.2	NA
4	2015	Arizona	NA	15.2	NA
5	2015	Arizona	NA	15.2	NA
6	2015	Arizona	NA	15.2	NA
7	2015	Arizona	3.2	15.2	-12.0
8	2015	Arizona	3.2	15.2	-12.0
9	2015	Arizona	3.1	15.2	-12.1
10	2015	Arizona	12.5	15.2	-2.66

```
# i 9,784 more rows
```

Use `n()` for sample size by group

There are other functions, such as `n()` count the number of observations.

```
yls %>%  
  group_by(YEAR) %>%  
  summarize(n = n(),  
            mean = mean(Data_Value, na.rm = TRUE))
```

```
# A tibble: 17 × 3  
  YEAR      n mean  
  <dbl> <int> <dbl>  
1  1999   372  26.1  
2  2000  1224  26.7  
3  2001   426  23.4  
4  2002  1016  25.2  
5  2003   498  21.3  
6  2004   611  20.7  
7  2005   636  21.8  
8  2006   518  21.8  
9  2007   516  20.0  
10 2008   483  18.2  
11 2009   686  18.3  
12 2010   447  17.8  
13 2011   521  17.8  
14 2012   244  15.5  
15 2013   685  16.7  
16 2014   334  15.7  
17 2015   577  15.2
```

Iterative summaries

Iterative summaries: **dplyr** `summarize()` and `across()` functions

Use the `across` function with `summarize()` to summarize across multiple columns of your data.

```
# General format - Not the code!  
across({ columns to go across }, ~ { summarization_function(.x, na.rm = ..) })
```


Iterative summaries: `dplyr` `summarize()` and `across()` functions

Use the `across` function with `summarize()` to summarize across multiple columns of your data.

```
# General format - Not the code!
```

```
across({ columns to go across }, ~ { summarization_function(.x, na.rm = ..) })
```

```
yts %>%
```

```
  group_by(YEAR) %>%
```

```
  summarize(across( c(Data_Value, Data_Value_Std_Err, Sample_Size),  
                    ~ mean(.x, na.rm = TRUE)))
```

```
# A tibble: 17 × 4
```

	YEAR	Data_Value	Data_Value_Std_Err	Sample_Size
	<dbl>	<dbl>	<dbl>	<dbl>
1	1999	26.1	1.98	1591.
2	2000	26.7	2.03	1743.
3	2001	23.4	1.79	2060.
4	2002	25.2	1.81	2653.
5	2003	21.3	1.92	2325.
6	2004	20.7	1.84	1246.
7	2005	21.8	2.17	1017.
8	2006	21.8	2.15	1191.
9	2007	20.0	1.96	1093.
10	2008	18.2	1.73	1203.
11	2009	18.3	1.90	1033.
12	2010	17.8	1.71	1202.
13	2011	17.8	1.84	1274.
14	2012	15.5	1.58	1052.

Iterative summaries: `dplyr` `summarize()` and `across()` functions

Another example using select helpers (??`tidyr_tidy_select`):

```
yts %>%  
  summarize(across( where(is.numeric), ~ mean(.x, na.rm = TRUE)))
```

A tibble: 1 × 7

	YEAR	Data_Value	Data_Value_Std_Err	Low_Confidence_Limit	High_Confidence_Limit	Sample_Size	Di
	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	
1	2006.	21.0	1.87	17.3	24.6	1505.	

Data Summarization on data frames

- Basic statistical summarization for numeric data
 - `rowMeans(x)`: takes the means of each row of x
 - `colMeans(x)`: takes the means of each column of x
 - `rowSums(x)`: takes the sum of each row of x
 - `colSums(x)`: takes the sum of each column of x
 - `summary(x)`: for data frames, displays the quantile information

summary() Function

Using `summary()` can give you rough snapshots of each numeric column (character columns are skipped):

```
summary(yts)
```

YEAR	LocationAbbr	LocationDesc	TopicType	TopicDesc
Min. :1999	Length:9794	Length:9794	Length:9794	Length:9794
1st Qu.:2002	Class :character	Class :character	Class :character	Class :character
Median :2006	Mode :character	Mode :character	Mode :character	Mode :character
Mean :2006				
3rd Qu.:2010				
Max. :2015				

MeasureDesc	DataSource	Response	Data_Value_Unit	Data_Value_Type
Length:9794	Length:9794	Length:9794	Length:9794	Length:9794
Class :character	Class :character	Class :character	Class :character	Class :character
Mode :character	Mode :character	Mode :character	Mode :character	Mode :character

Data_Value	Data_Value_Footnote_Symbol	Data_Value_Footnote	Data_Value_Std_Err	Low_Confidence
Min. : 0.00	Length:9794	Length:9794	Min. : 0.000	Min. : 0.000
1st Qu.: 3.20	Class :character	Class :character	1st Qu.: 0.600	1st Qu.: 1.900
Median :11.30	Mode :character	Mode :character	Median : 1.300	Median : 8.500
Mean :20.97			Mean : 1.874	Mean :17.300
3rd Qu.:39.10			3rd Qu.: 2.500	3rd Qu.:31.600
Max :98.00			Max :16.100	Max :97.000

Summary

- summary stats (`mean()`) work with `pull()`
- `count(x)`: what unique values do you have?
 - `pull()` to get vectors
 - `unique()` combined with `length()`
- `group_by()`: changes all subsequent functions
 - combine with `summarize()` to get statistics per group
 - combine with `across()` to programmatically select columns
- `summary(x)`: quantile information