

Merging Data Sets

Data Wrangling in R

Joining in `dplyr`

- Merging/joining data sets together - usually on key variables, usually "id"
- `?join` - see different types of joining for `dplyr`
- `inner_join(x, y)` - only rows that match for `x` and `y` are kept
- `full_join(x, y)` - all rows of `x` and `y` are kept
- `left_join(x, y)` - all rows of `x` are kept even if not merged with `y`
- `right_join(x, y)` - all rows of `y` are kept even if not merged with `x`
- `anti_join(x, y)` - all rows from `x` not in `y` keeping just columns from `x`.

Merging: Simple Data

base has baseline data for ids 1 to 10 and Age

```
base <- tibble(id = 1:10, Age = seq(55, 60, length=10))  
head(base, 2)
```

```
# A tibble: 2 x 2  
      id   Age  
  <int> <dbl>  
1     1  55  
2     2 55.6
```


visits has ids 2 to 11, 3 different visits, and an outcome

```
visits <- tibble(id = rep(2:11, 3), visit= rep(1:3, 10),  
                 Outcome = seq(10, 50, length=30))  
head(visits, 2)
```

```
# A tibble: 2 x 3  
      id visit Outcome  
  <int> <int>   <dbl>  
1     2     1     10  
2     3     2    11.4
```

Inner Join

`inner_join(x, y)`



<https://github.com/gadenbuie/tidyexplain/blob/master/images/inner-join.gif>

Inner Join

```
ij = inner_join(base, visits)
```

```
Joining, by = "id"
```

```
dim(ij)
```


```
[1] 27  4
```

```
head(ij)
```

```
# A tibble: 6 x 4
   id    Age visit Outcome
<int> <dbl> <int>   <dbl>
1     2  55.6     1      10
2     2  55.6     2     23.8
3     2  55.6     3     37.6
4     3  56.1     2     11.4
5     3  56.1     3     25.2
6     3  56.1     1     39.0
```

Left Join

`left_join(x, y)`



<https://github.com/gadenbuie/tidyexplain/blob/master/images/left-join.gif>

Left Join

```
lj = left_join(base, visits)
```

```
Joining, by = "id"
```

```
dim(lj)
```

```
[1] 28  4
```

```
head(lj)
```

```
# A tibble: 6 x 4
   id    Age visit Outcome
<int> <dbl> <int>   <dbl>
1     1    55     NA      NA
2     2   55.6     1     10
3     2   55.6     2    23.8
4     2   55.6     3    37.6
5     3   56.1     2    11.4
6     3   56.1     3    25.2
```

Install tidylog package to log outputs

```
# install.packages("tidylog")
library(tidylog)
left_join(base, visits)
```

Joining, by = "id"

left_join: added 2 columns (visit, Outcome)

```
> rows only in x      1
```

```
> rows only in y    ( 3)
```

```
> matched rows      27      (includes duplicates)
```

```
>                      =====
```


```
> rows total        28
```

A tibble: 28 x 4

	id	Age	visit	Outcome
	<int>	<dbl>	<int>	<dbl>
1	1	55	NA	NA
2	2	55.6	1	10
3	2	55.6	2	23.8
4	2	55.6	3	37.6
5	3	56.1	2	11.4
6	3	56.1	3	25.2

Right Join

`right_join(x, y)`



<https://github.com/gadenbuie/tidyexplain/blob/master/images/right-join.gif>

Right Join

```
rj = right_join(base, visits)
```

```
Joining, by = "id"
```

```
right_join: added 2 columns (visit, Outcome)
```

```
> rows only in x ( 1)
```

```
> rows only in y    3
```

```
> matched rows      27
```

```
>                      ====
```

```
> rows total        30
```

Left Join: Switching arguments

```
lj2 = left_join(visits, base)
```

```
Joining, by = "id"
```

```
left_join: added one column (Age)
```

```
> rows only in x      3
```

```
> rows only in y    ( 1)
```


```
> matched rows      27
```

```
>                      =====
```

```
> rows total        30
```

Full Join

`full_join(x, y)`



<https://github.com/gadenbuie/tidyexplain/blob/master/images/full-join.gif>

Full Join

```
fj = full_join(base, visits)
```

Joining, by = "id"

full_join: added 2 columns (visit, Outcome)

```
> rows only in x      1
```

```
> rows only in y      3
```

```
> matched rows       27      (includes duplicates)
```

```
>                      =====
```

```
> rows total          31
```

Full Join

Note what tidylog means by `includes` duplicates. Data from `base` is being duplicated.

```
# fj = full_join(base, visits)
head(fj, 10)
```

```
# A tibble: 10 x 4
      id    Age visit Outcome
  <int> <dbl> <int>   <dbl>
1     1    55    NA      NA
2     2   55.6     1     10
3     2   55.6     2    23.8
4     2   55.6     3    37.6
5     3   56.1     2    11.4
6     3   56.1     3    25.2
7     3   56.1     1    39.0
8     4   56.7     3    12.8
9     4   56.7     1    26.6
10    4   56.7     2    40.3
```

Duplicated

The `duplicated` command can give you indications if there are duplications in a **vector**:

```
duplicated(1:5)
```

```
[1] FALSE FALSE FALSE FALSE FALSE
```

```
duplicated(c(1:5, 1))
```

```
[1] FALSE FALSE FALSE FALSE FALSE TRUE
```

```
fj %>% mutate(dup_id = duplicated(id))
```

```
# A tibble: 31 x 5
   id   Age visit Outcome dup_id
  <int> <dbl> <int>   <dbl> <lgl>
1     1   55    NA     NA  FALSE
2     2  55.6     1    10  FALSE
3     2  55.6     2   23.8  TRUE
4     2  55.6     3   37.6  TRUE
5     3  56.1     2   11.4  FALSE
6     3  56.1     3   25.2  TRUE
7     3  56.1     1   39.0  TRUE
8     4  56.7     3   12.8  FALSE
9     4  56.7     1   26.6  TRUE
10    4  56.7     2   40.3  TRUE
# ... with 21 more rows
```

Using the **by** argument

By default - uses intersection of column names. If `by` specified, then uses that.

```
# for multiple, by = c(col1, col2)
head(full_join(base, visits, by = "id"))
```

```
# A tibble: 6 x 4
   id    Age visit Outcome
<int> <dbl> <int>   <dbl>
1     1    55     NA      NA
2     2   55.6     1     10
3     2   55.6     2    23.8
4     2   55.6     3    37.6
5     3   56.1     2    11.4
6     3   56.1     3    25.2
```


Using the **by** argument

You can use `by` if the column names don't match exactly.

```
base2 = base %>% rename(patient = id) # rename the column  
head(full_join(base2, visits, by = c("patient" = "id")))
```

```
# A tibble: 6 x 4  
  patient    Age visit Outcome  
  <int> <dbl> <int>   <dbl>  
1         1    55     NA     NA  
2         2   55.6      1    10  
3         2   55.6      2   23.8  
4         2   55.6      3   37.6  
5         3   56.1      2   11.4  
6         3   56.1      3   25.2
```